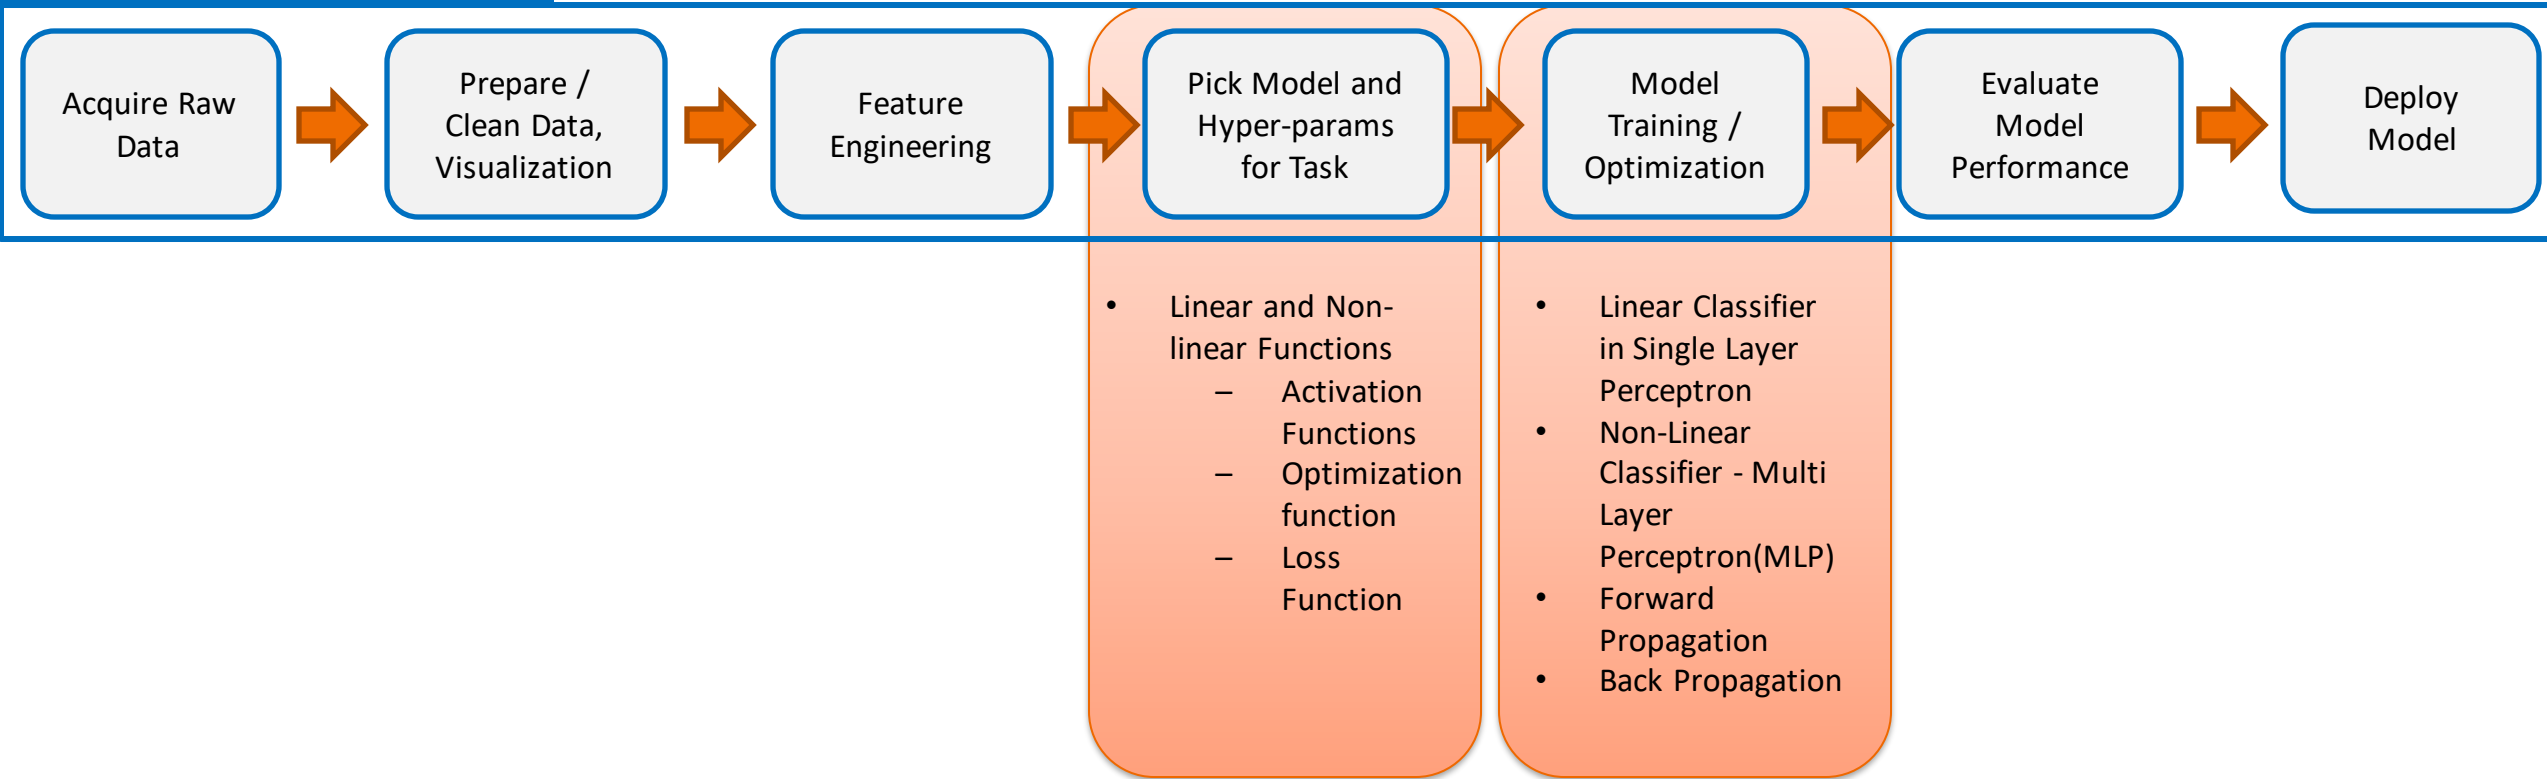
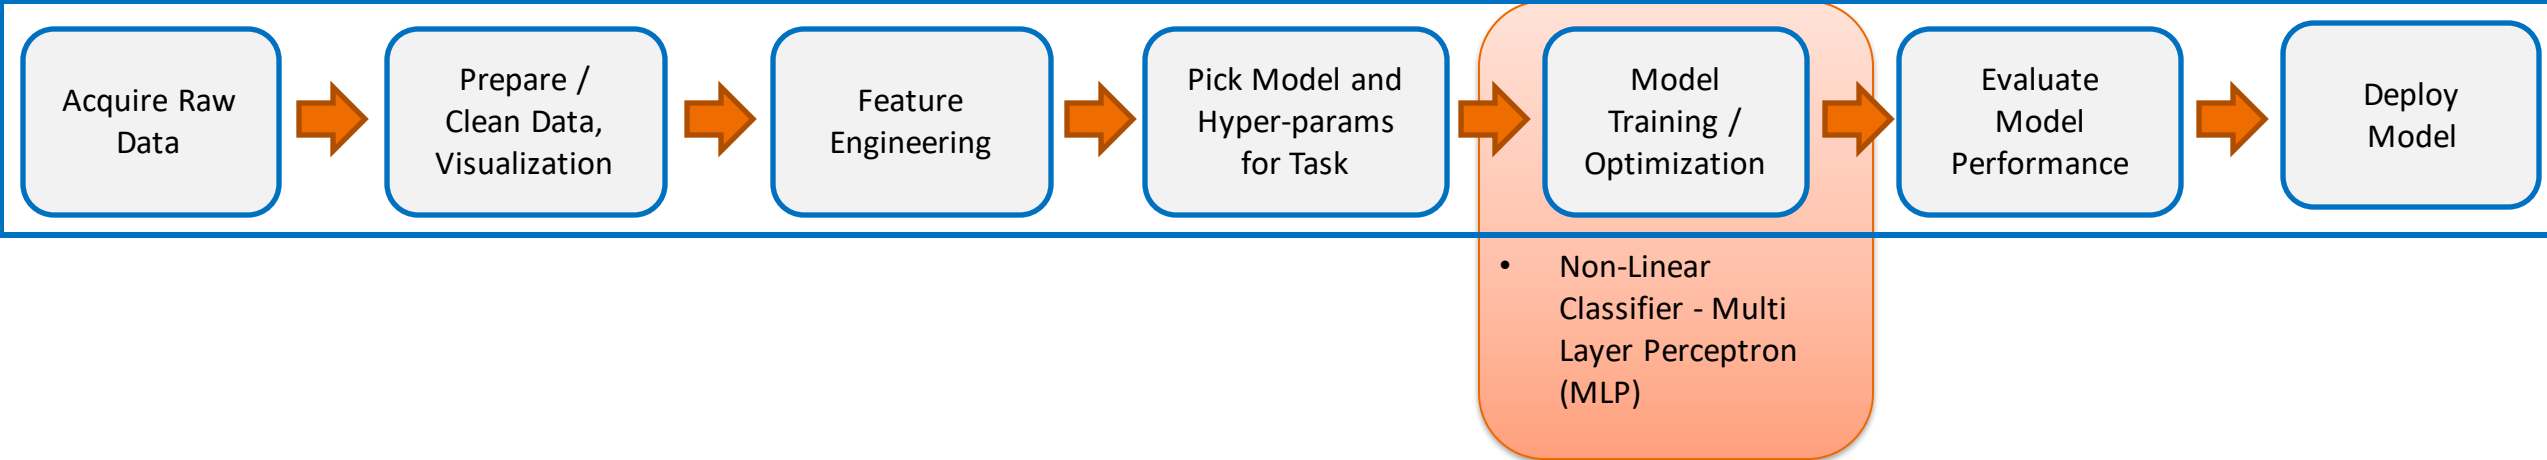


Focus for this lecture

Task → Classification

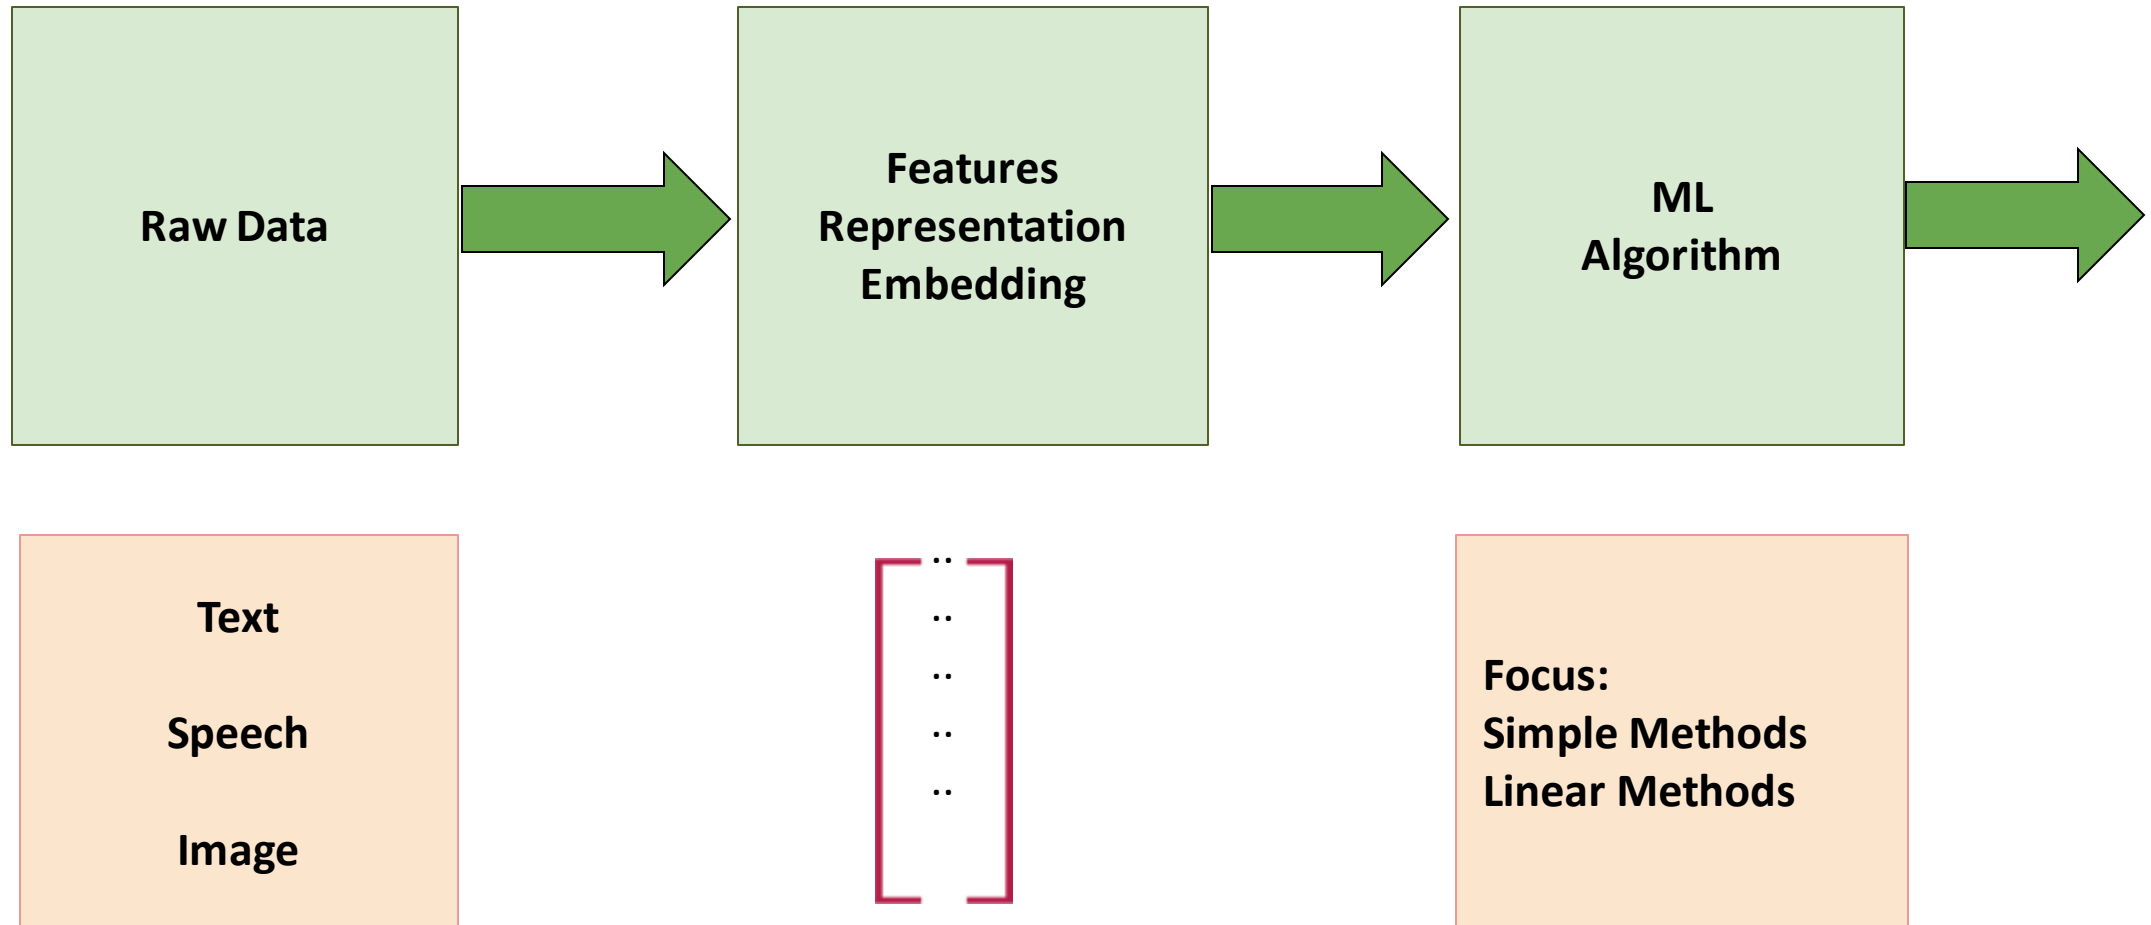




Non-Linear Classification

MLP

Pipeline




Spectrum of Classifiers




Linear and Non Linear Functions


$$W^T X = w_0 + w_1 x_1 = x_1 + 2$$




$$\sinh(W^T X) = \sinh(w_1 x_1) = \sinh(0.5 x_1)$$




$$W^T X^2 = w_1 x_1^2 = x_1^2$$





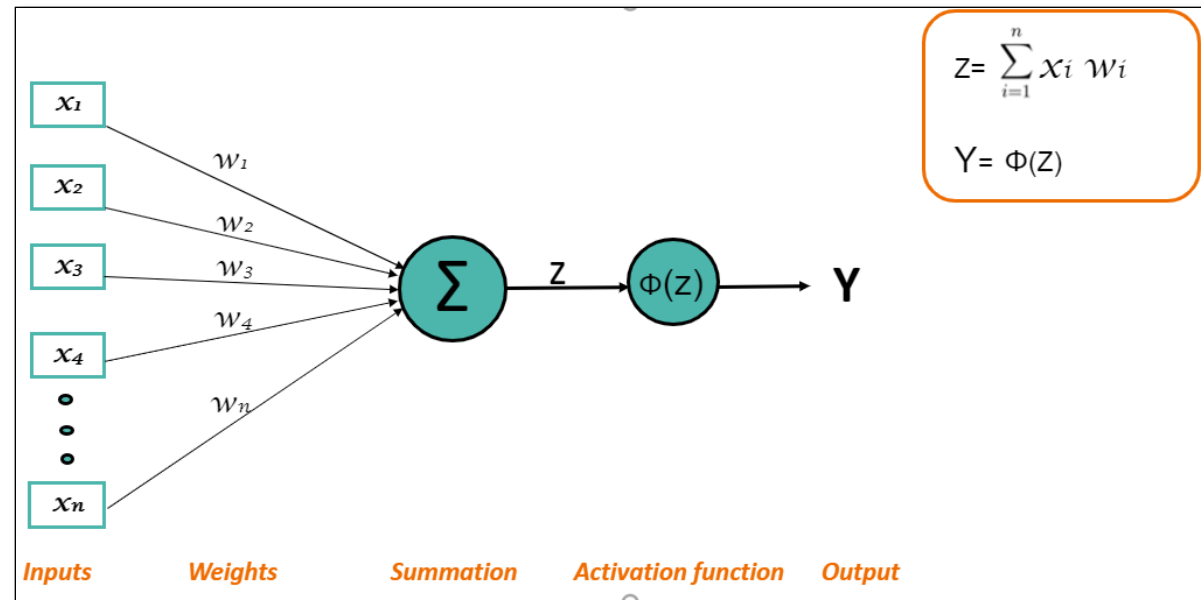
$$e^{-W^T X} = e^{(w_1 x_1)} = e^{(-2 x_1)}$$



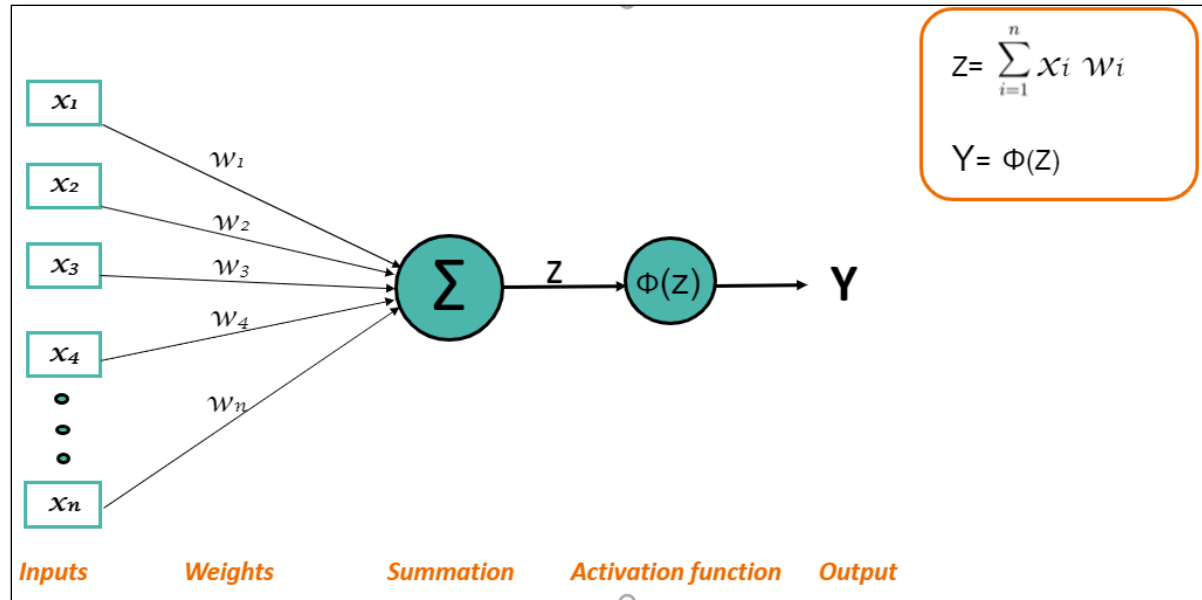
Single Layer Perceptron

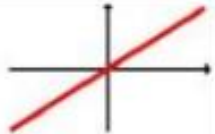
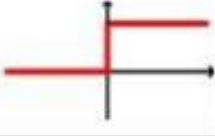
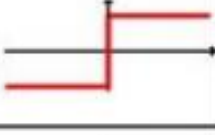


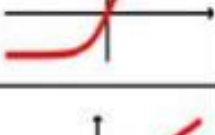



Why Use Only One Neuron ?




Activation Functions




Activation Function	Equation	Example	1D Graph
Linear	$\phi(z) = z$	Adaline, linear regression	
Unit Step (Heaviside Function)	$\phi(z) = \begin{cases} 0 & z < 0 \\ 0.5 & z = 0 \\ 1 & z > 0 \end{cases}$	Perceptron variant	
Sign (signum)	$\phi(z) = \begin{cases} -1 & z < 0 \\ 0 & z = 0 \\ 1 & z > 0 \end{cases}$	Perceptron variant	
Piece-wise Linear	$\phi(z) = \begin{cases} 0 & z \leq -\frac{1}{2} \\ z + \frac{1}{2} & -\frac{1}{2} \leq z \leq \frac{1}{2} \\ 1 & z \geq \frac{1}{2} \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression,	
Hyperbolic Tangent (tanh)	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$		
ReLU	$\phi(z) = \begin{cases} 0 & z < 0 \\ z & z > 0 \end{cases}$		

Deep Neural Networks (Multi Layer Perceptron (MLP))



More layers : Able to model complex decision boundaries



Note: Every neuron has a nonlinearity Φ inside. This is required to model non-linear boundaries

Multi layer Perceptron

—— Popular Artificial Neural Networks ——

Simple Problem

- Given #Bedrooms, Area (Sqft), LocalityIndex (#houses in the neighbourhood),
 - Predict whether the price of the house is “high” (1) or “low” (0).


Eg. House price from attributes

Bedrooms	Sq. Feet	Neighborhood (no. of houses in the locality)	Price high or low? High (1), Low (0)
3	2000	90	1
2	800	143	0
2	850	167	0
1	550	267	0
4	2000	396	1


MLP Architecture

- MLP consists of at least three layers
(input layer + hidden layer(s) + output layer)
- Each layer (except input layer) consists of neurons that use an activation function (eg. step or sigmoid)
- Learning/Training: Backpropagation is used (not for today!!)

MLP




A Note on Notations

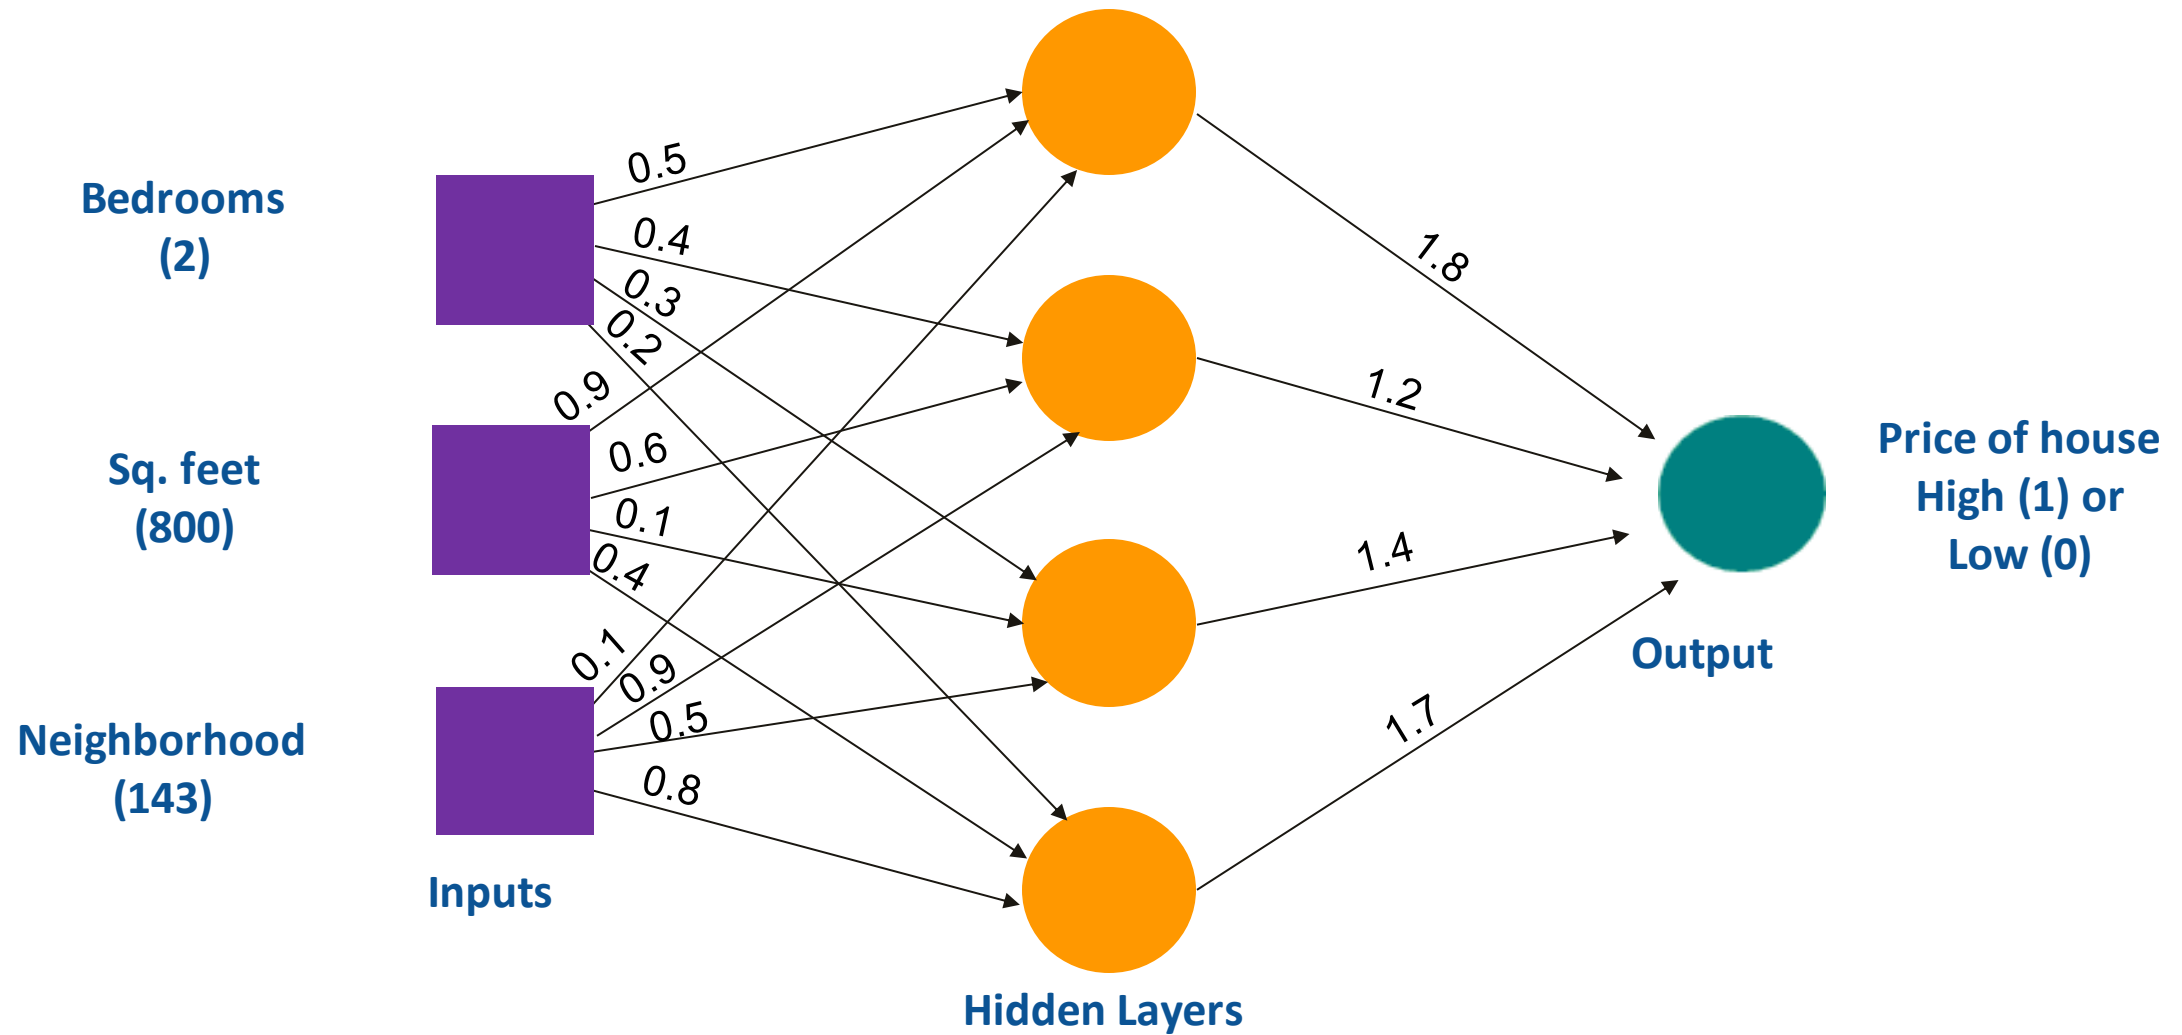


$$\begin{bmatrix} z_1 \\ z_2 \\ z_3 \\ z_4 \end{bmatrix} = \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \\ w_{41} & w_{42} & w_{43} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$


$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix} = \begin{bmatrix} \phi(z_1) \\ \phi(z_2) \\ \phi(z_3) \\ \phi(z_4) \end{bmatrix}$$




Initialize weights




Weights at the first neuron




Weights at the second neuron




Weights at the third neuron



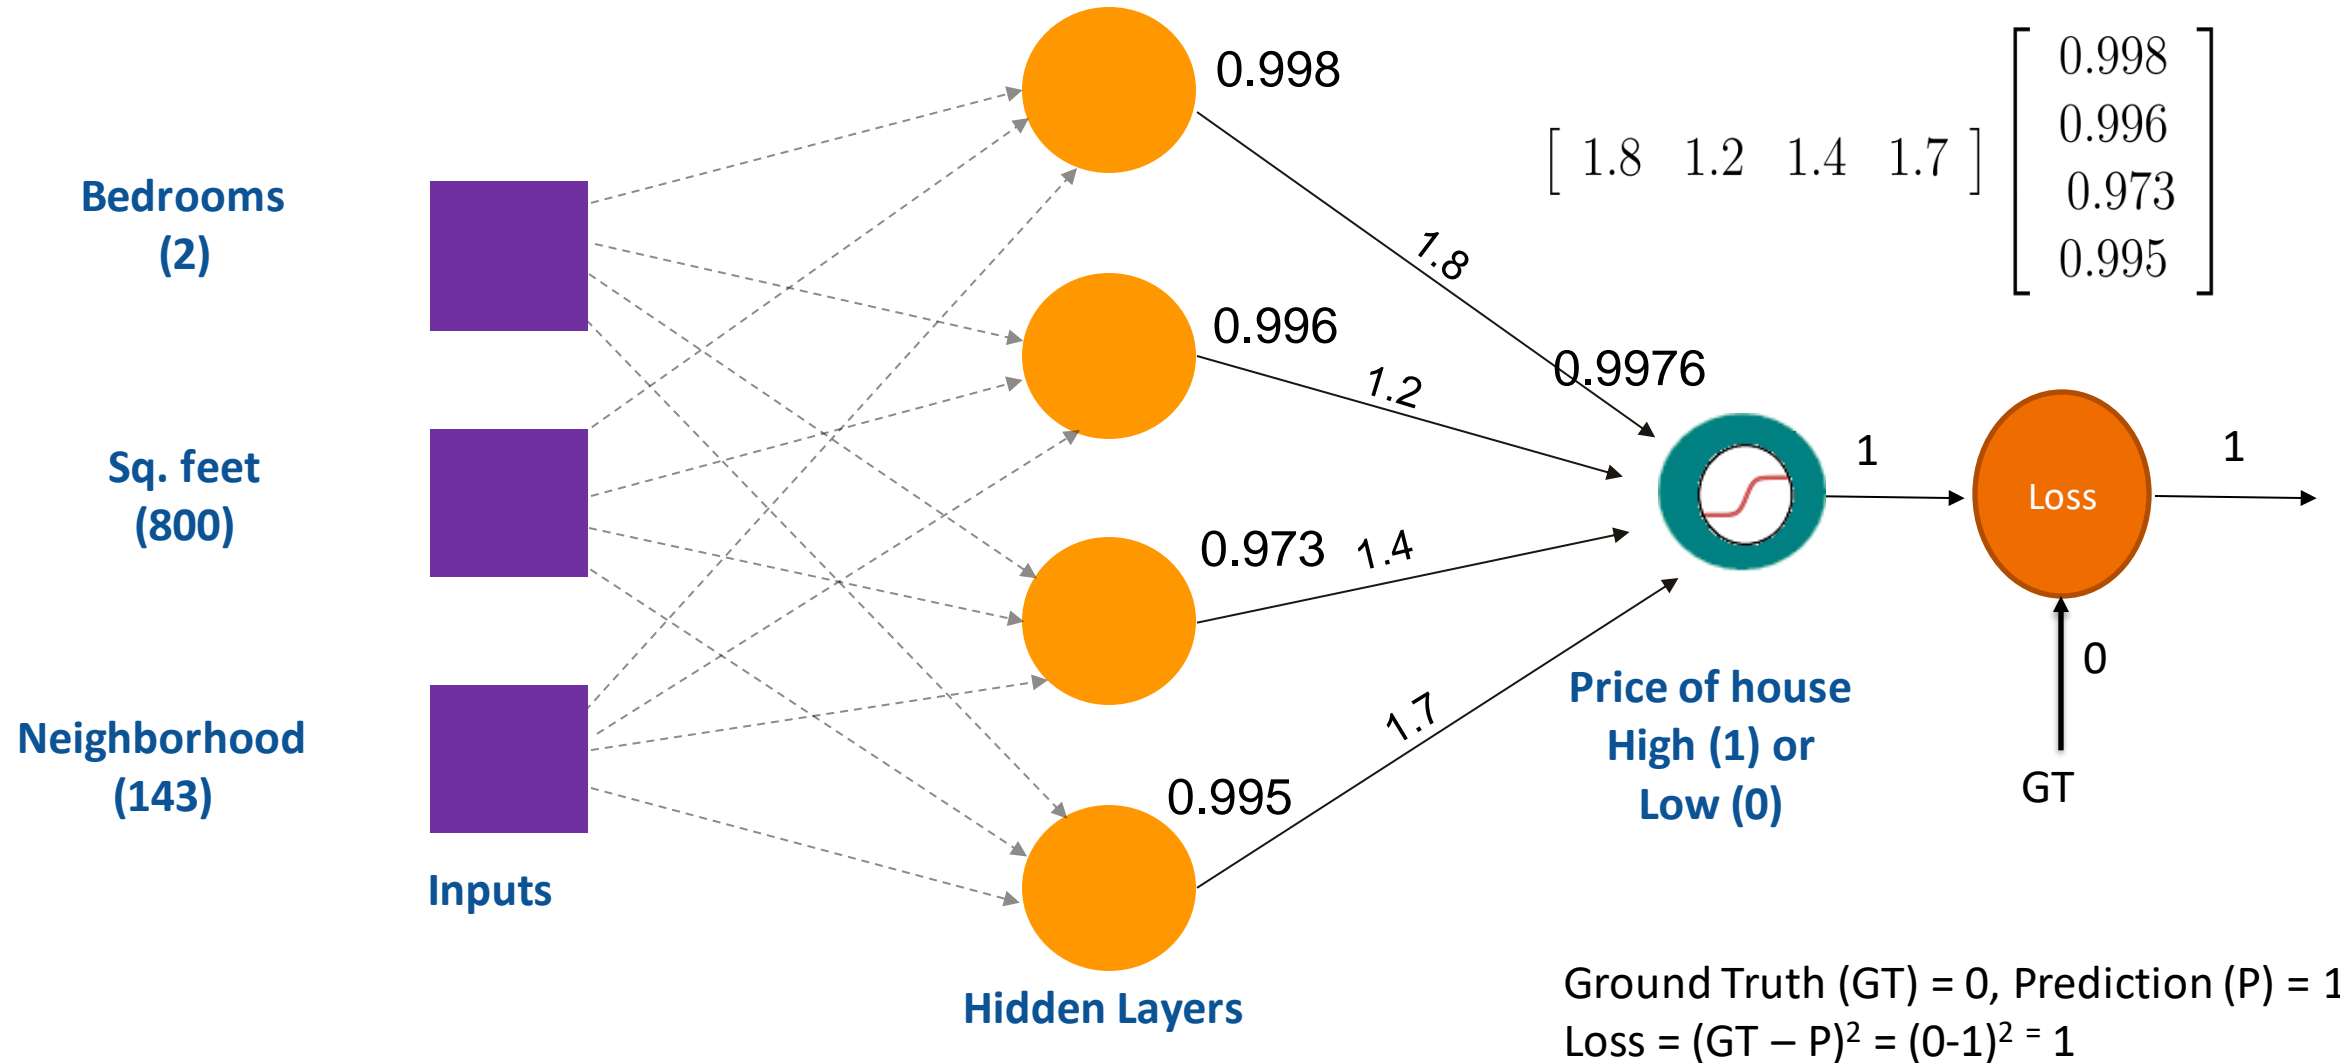
Weights at the fourth neuron



Inputs to the next layer



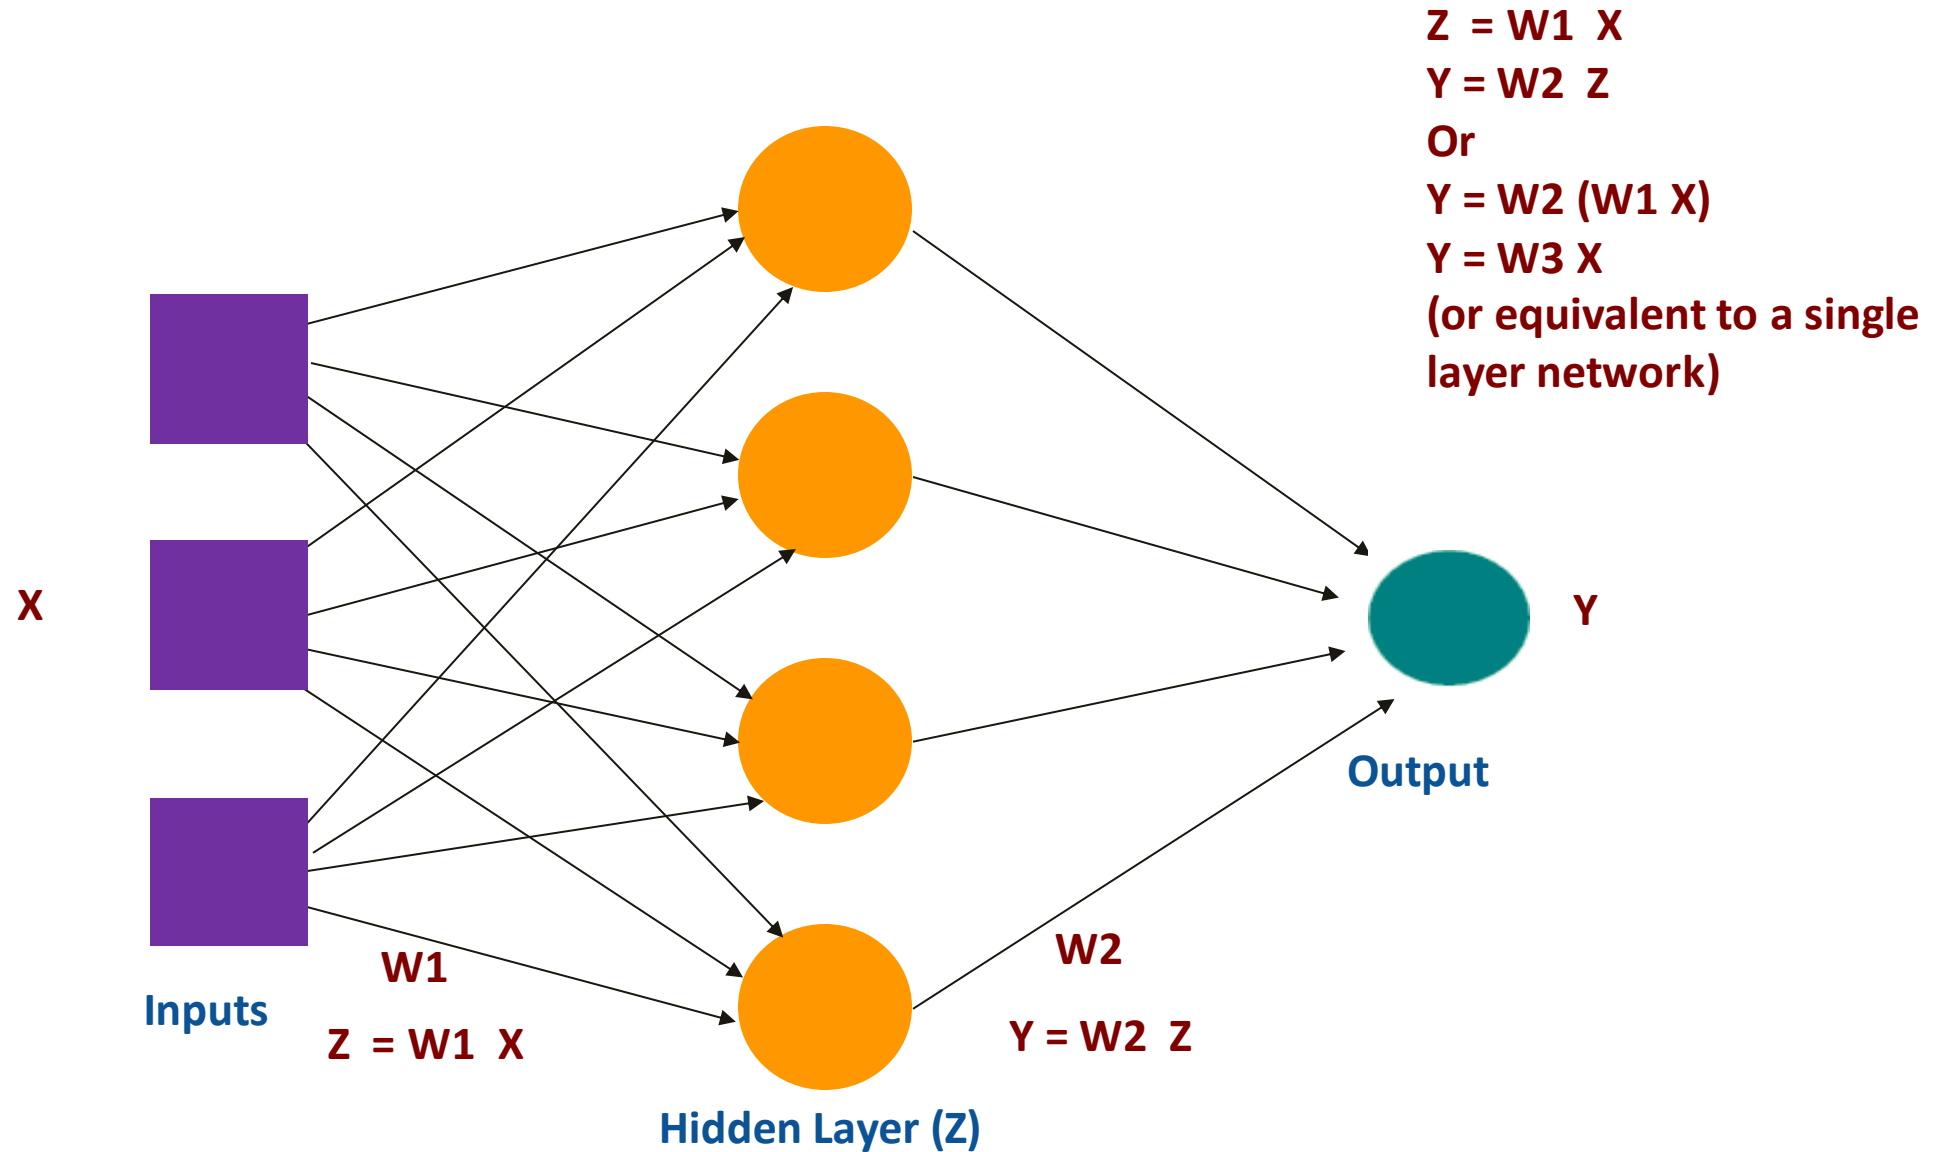
Computations in next layer



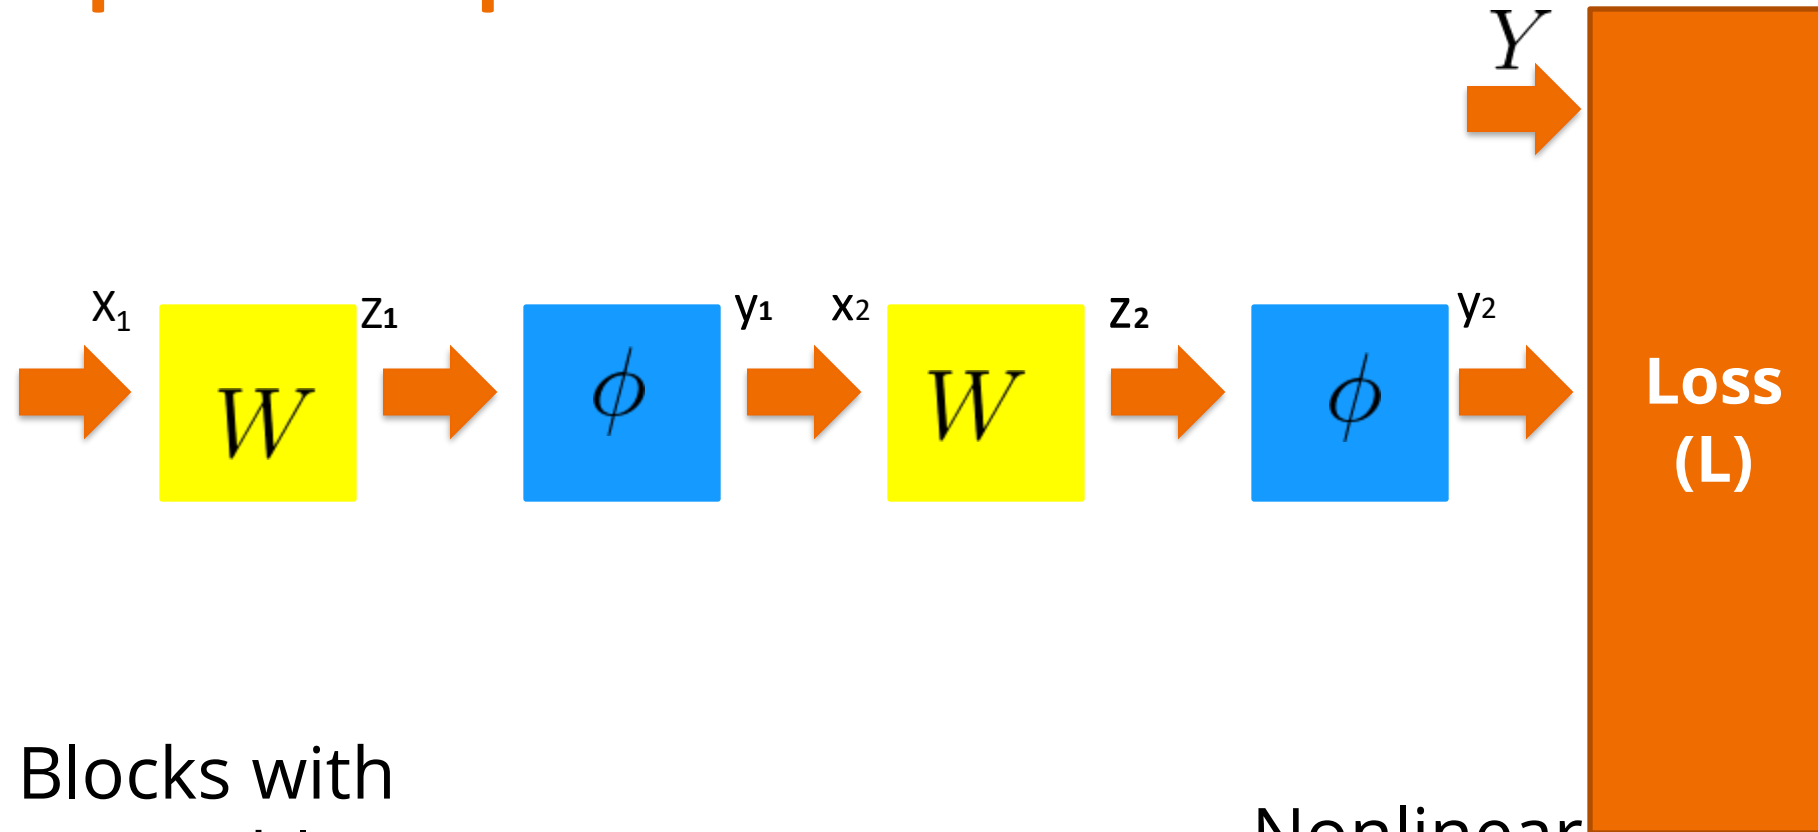
Why Nonlinearity in MLP?

Answer?

Comment: Limitation of “Linear MLP”




A simpler view point



Blocks with
Learnable
parameters
Matrix
Multiplication

Nonlinear
functions
(often non
learnable)

Loss or Objective




Objective: Find out the best parameters which will minimize the loss.

$$W^* = \arg \min_W \sum_{i=1}^N L(y'_i, y_i; W) \longrightarrow \text{Weight Vector}$$

$$L = \frac{1}{2} \|y'_i - y_i\|_2^2 \longrightarrow \text{E.g. Squared Loss}$$

Back Propagation



Solution: Iteratively update W along the direction where loss decreases.

Each layer's weights are updated based on the derivative of its output w.r.t. input and weights

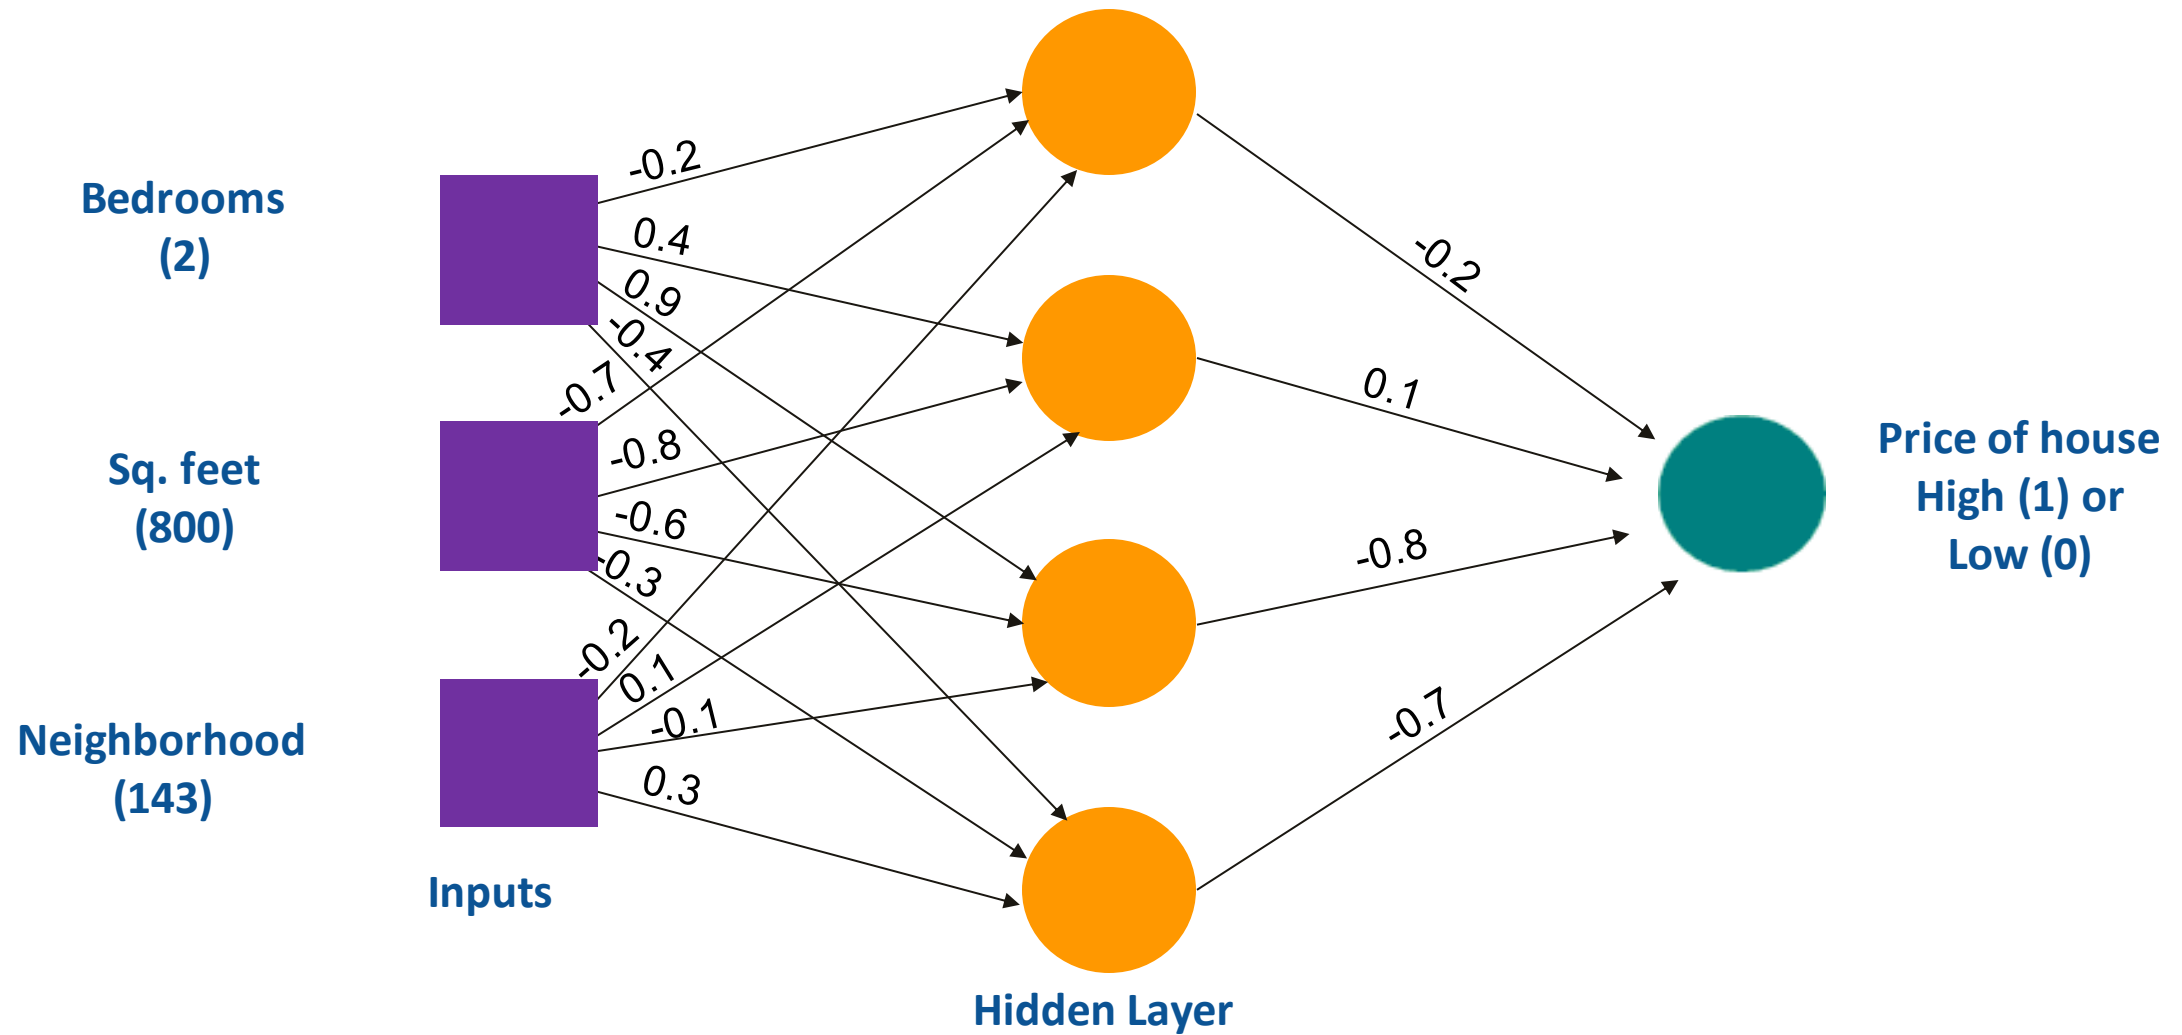
Backpropagation Algorithm

- **Input:** A set of examples (x_i, y_i)
- **Objective:** Update weights **W** so that error is small/optimal.
- **How:** For the set (or subset) of examples,
 1. Compute output
 2. Update weights (gradient descent)
 3. Repeat 1-3 until convergence.


Loss or Error

- $0.998 * 1.8 + 0.996 * 1.2 + 0.973 * 1.4 + 0.995 * 1.7$
 $= \phi(6.04) = 0.9976$
- For threshold nonlinearity Output = 1
- The actual class (0) deviates from the predicted class (1)
- The squared error or Loss is $(1-0) * (1-0) = 1$
- The weights to be updated by **backpropagation to reduce the error**


Weights obtained by backpropagation




Weights at the first neuron




Weights at the second neuron



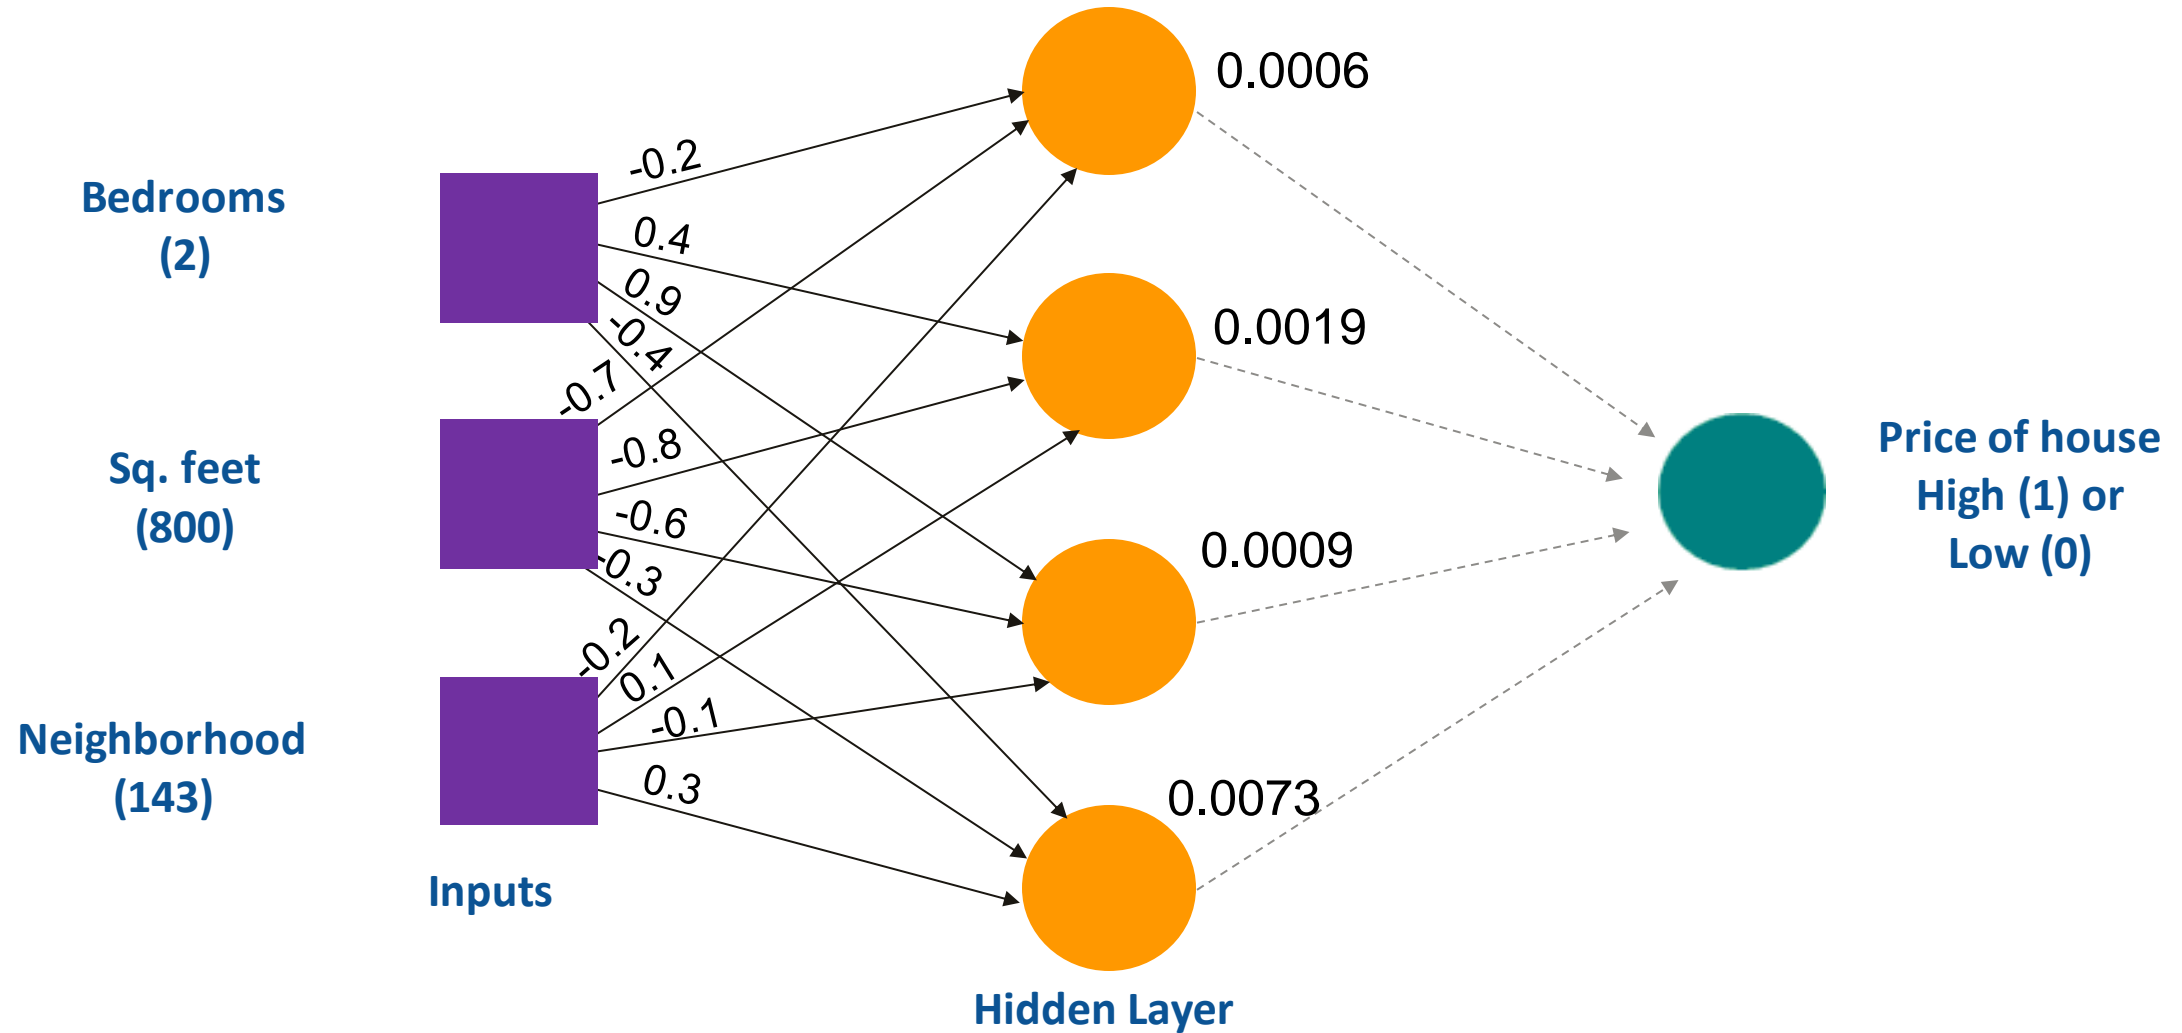
Weights at the third neuron



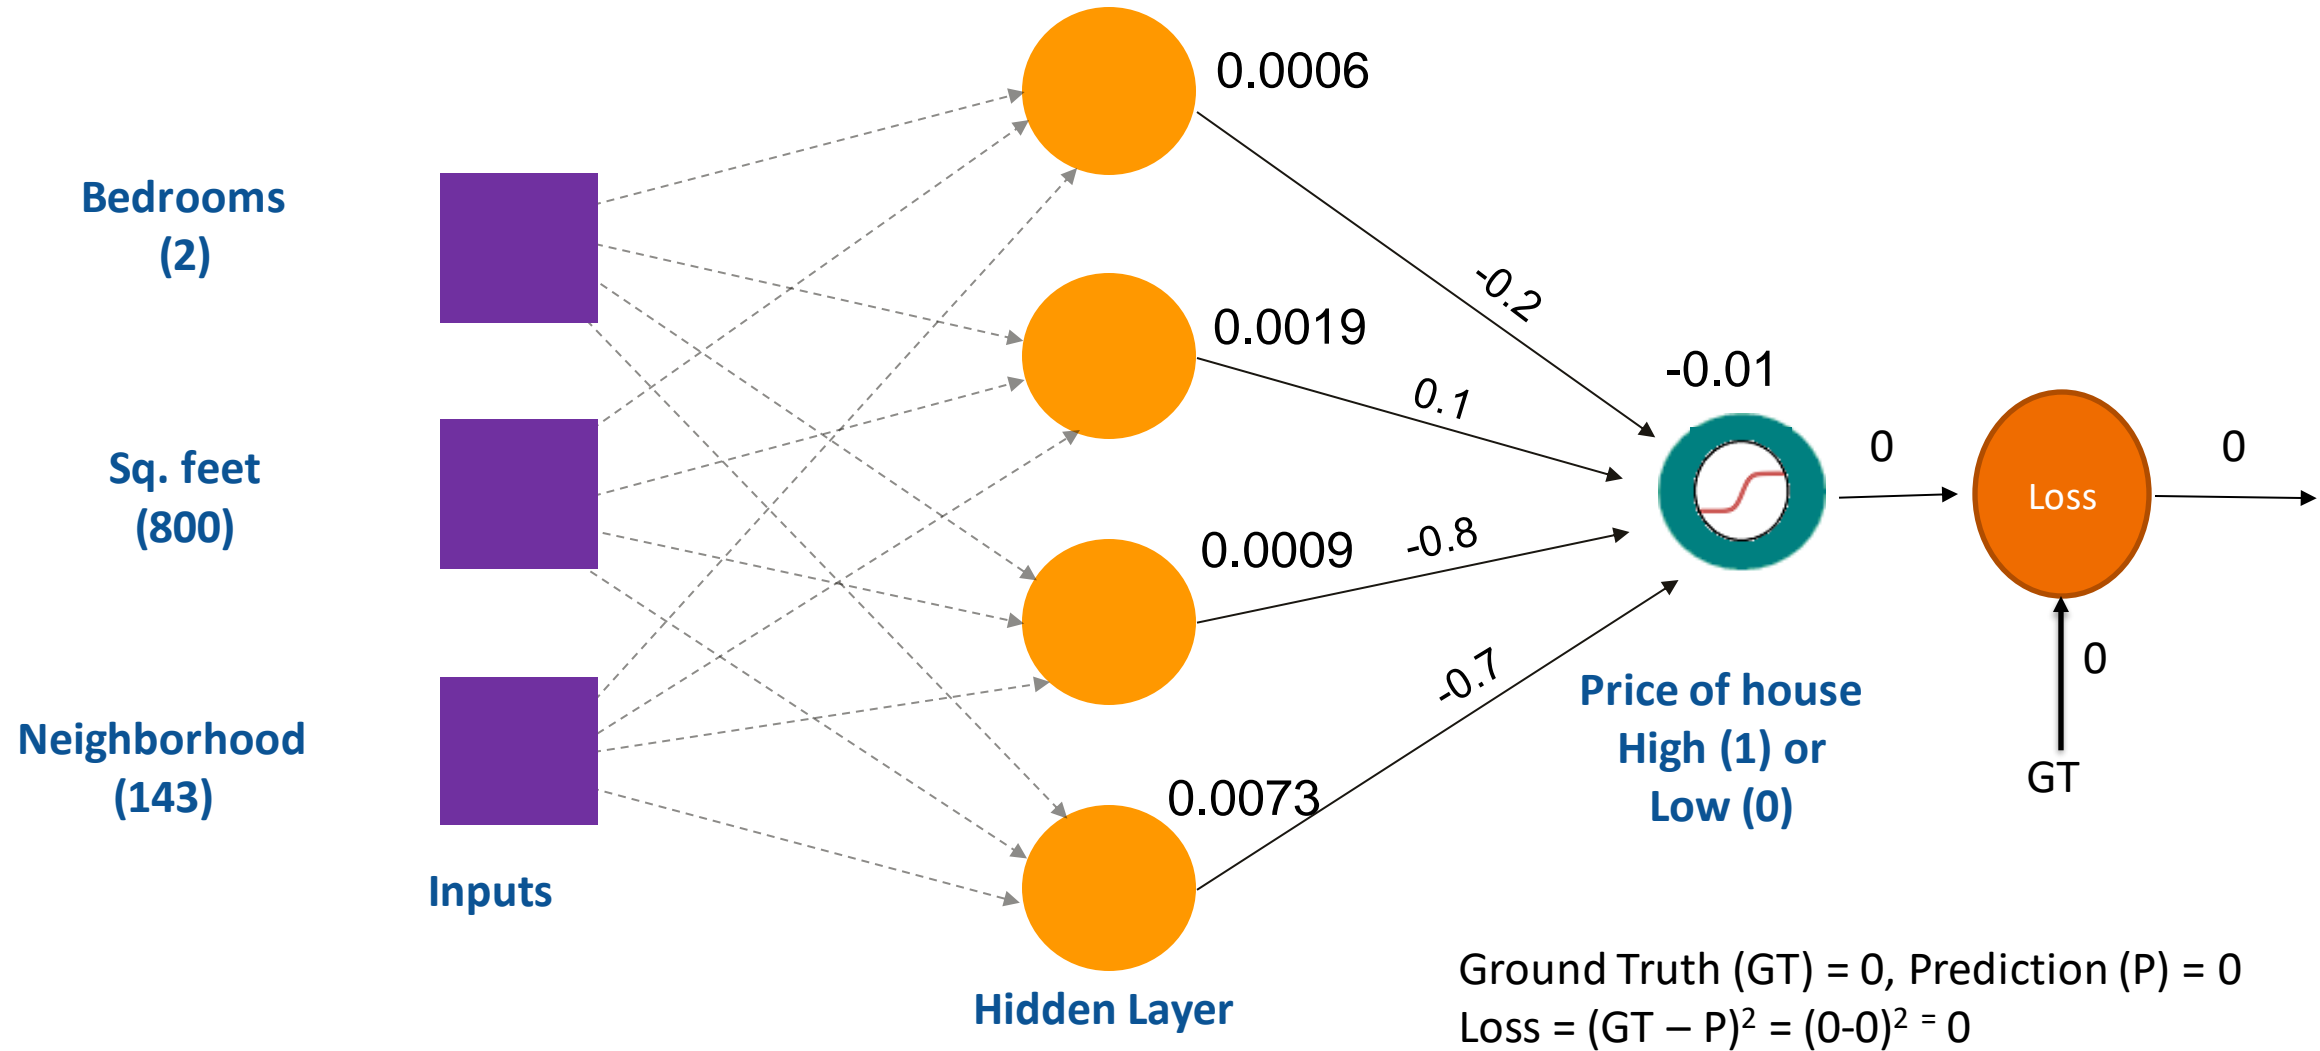
Weights at the fourth neuron



Activation function applied at first layer



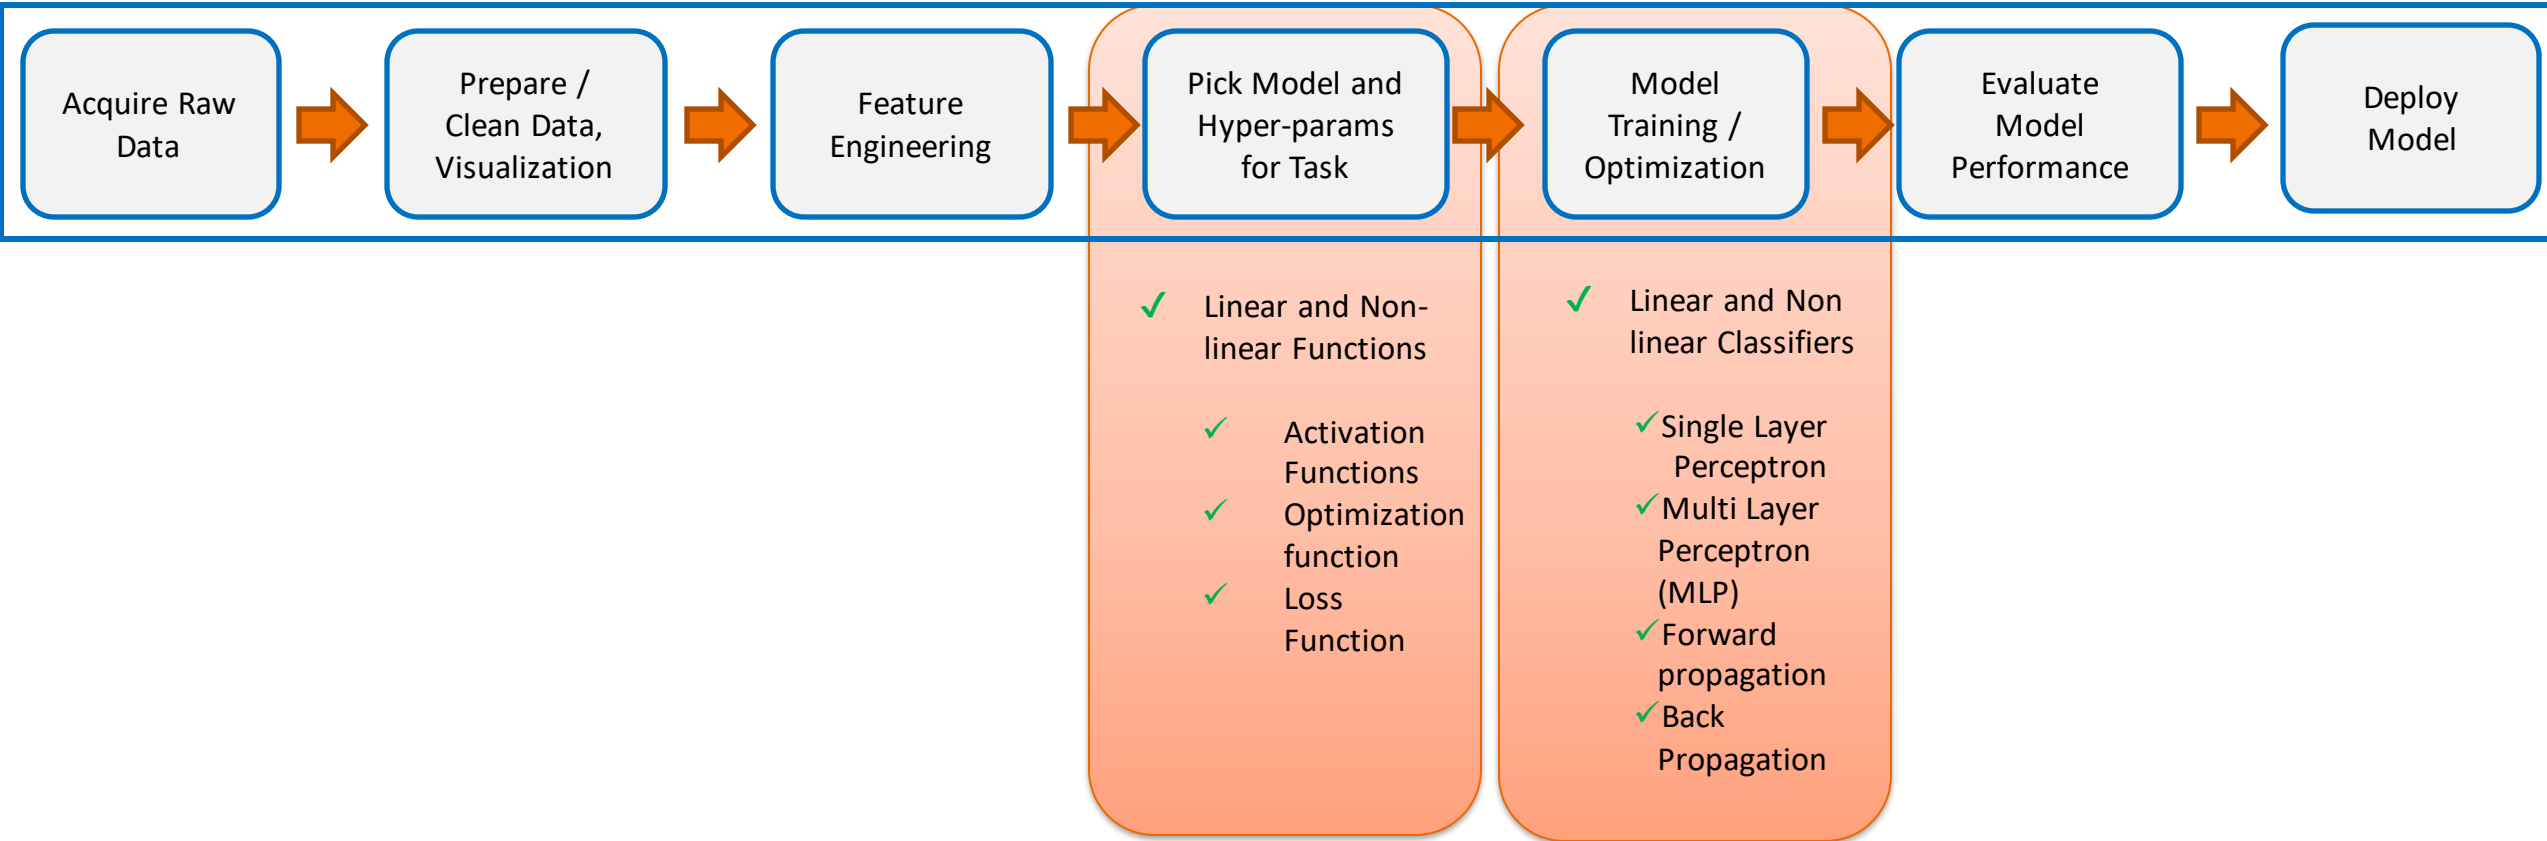
Activation function applied at Second layer



Loss/Error at this stage

- $0.0006 * -0.2 + 0.0019 * 0.1 + 0.0009 * -0.8 + 0.0073 * -0.7$
 $= -0.001$
- The square error or Loss with activation is $(0-0)*(0-0) = 0$
- The actual class (0) and the predicted class is also (0)
- Loss is now reduced from 1 to 0; appropriate weights are also found!!

Journey so far...



Case Study

Classification and Regression using MLP



Eg. House price from attributes (Classification)

Bedrooms	Sq. Feet	Neighborhood (no. of houses in the locality)	Price high or low? High (1), Low (0)
3	2000	90	1
2	800	143	0
2	850	167	0
1	550	267	0
4	2000	396	1

MLP for Classification

$$W1 = \begin{bmatrix} 0.2889 & 0.4082 & 0.2960 \\ -0.4042 & 0.2009 & -0.4359 \\ 0.3905 & -0.5748 & -0.0342 \\ -0.2267 & -0.1247 & -0.2257 \end{bmatrix}$$

$$W2 = [-0.5873, -0.0893, -0.1946, 0.0869]$$





Eg. House price from attributes (Regression)

Bedrooms	Sq. Feet	Neighborhood (no. of houses in the locality)	Price (in lakhs)
3	2000	90	23.0
2	800	143	8.0
2	850	167	9.0
1	550	267	9.0
4	2000	396	25.0

MLP for Regression

$$W1 = \begin{bmatrix} 0.210 & -0.079 & 0.0455 \\ -0.0356 & -0.835 & -0.0344 \\ -0.0027 & 0.651 & -0.0473 \\ 0.648 & 0.301 & 0.234 \end{bmatrix}$$

$$W2 = [2.4267, -1.8249, 0.4363, 2.0402]$$



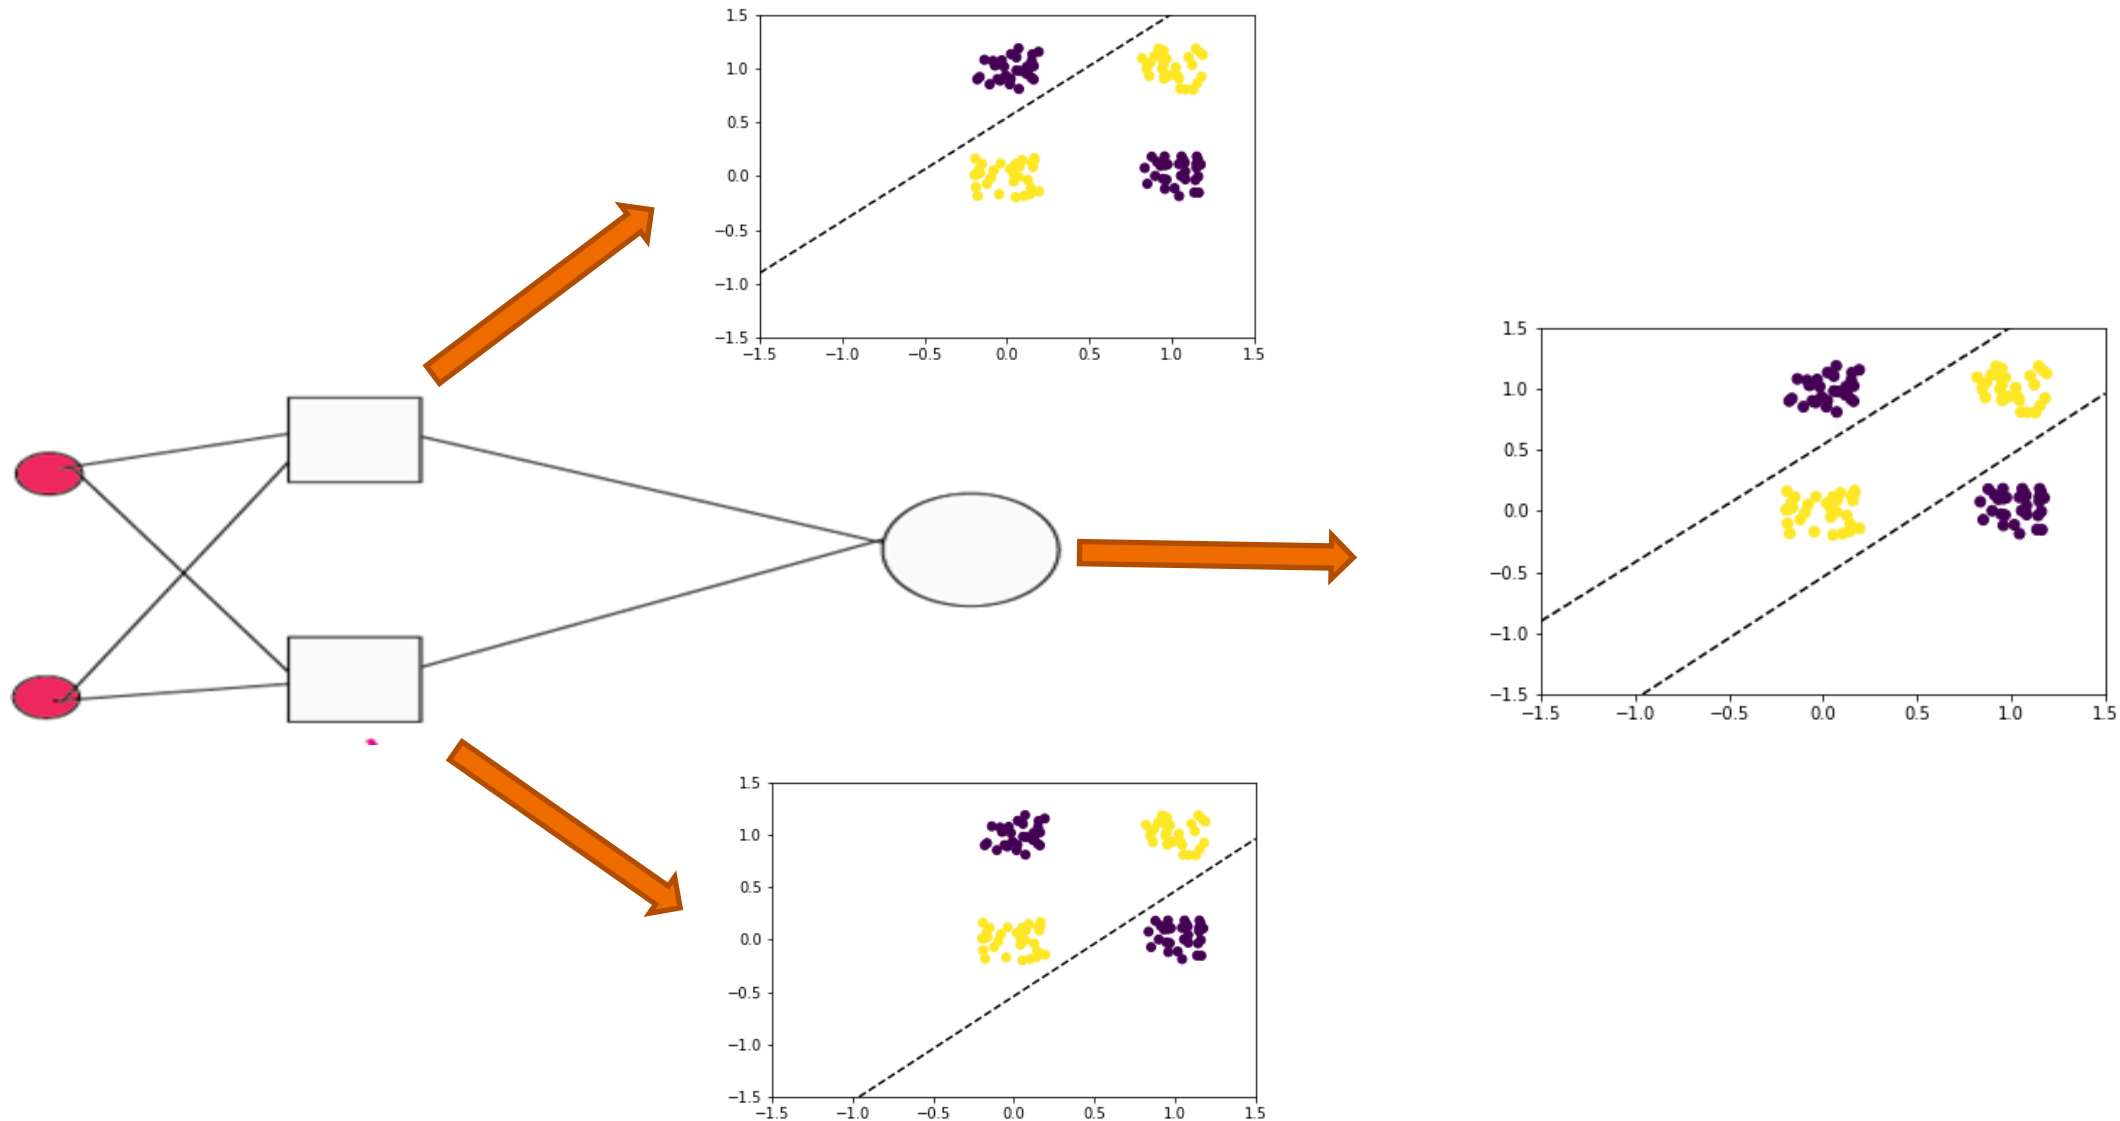
Questions?

Intuitive Explanation







—

—

How MLP Works? (A naïve view)











Deeper Networks







Case Study

What layers and neurons do in MLP?

What do neurons do?






What do layers do?



1st layer draws linear boundaries

2nd layer combines the boundaries


3rd layer can generate arbitrarily complex boundaries



Questions?


Activation Functions/Nonlinearities

Sigmoid




$$y = \frac{1}{1 + e^{-x}}$$

tanh




$$y = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

ReLU



$$y = \max(0, x)$$

Leaky ReLU




$$y = \begin{cases} x & \text{if } x > 0 \\ 0.01x & \text{otherwise} \end{cases}$$

maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$


Loss Functions



Mean Squared Loss

$$L(y, y') = (y - y')^2$$


y - Actual value
 y' - Predicted value




Mean Absolute Loss

$$L(y, y') = |y - y'|$$

Loss Functions



Cross Entropy Loss




Hinge Loss

y - Actual value
 y' - Predicted value

$$L(y, y') = -(y \log(y') + (1 - y) \log(1 - y'))$$

$$L(y, y') = \max(0, 1 - y * y')$$


Loss Functions (Regression)



Mean Squared Loss

$$L(y, y') = (y - y')^2$$

y - Actual value
 y' - Predicted value




Mean Absolute Loss

$$L(y, y') = |y - y'|$$

Loss Functions (Regression)


Mean Squared Loss
(Blue)



Mean Absolute Loss
(Red)

$$L(y, y') = (y - y')^2 \quad \begin{array}{l} y - \text{Actual value} \\ y' - \text{Predicted value} \end{array} \quad L(y, y') = |y - y'|$$

Loss Functions (Classification)



Cross Entropy Loss

$$L(y, y') = -(y \log(y') + (1 - y) \log(1 - y'))$$

y - Actual value

y' - Predicted value


Multi Class Classification using MLP

Input: (x_i, y_i)

$$x = [x_1, x_2, x_3, \dots, x_n]$$

Encode label y as


- $[1,0,0]$ for class 1
- $[0,1,0]$ for class 2
- $[0,0,1]$ for class 3



Softmax

- Normalizes the output.
- K is total number of classes

$$z_n = \frac{e^{x_n}}{\sum_{i=1}^K e^{x_i}}$$



```
Out[12]: array([ 6.,  0.,  5.,  3.,  8.])
```

```
In [8]: exp = (np.e)**(x)
        exp
```

executed in 6ms, finished 01:47:23 2018-08-21

```
Out[8]: array([ 4.03428793e+02,  1.00000000e+00,  1.48413159e+02,
                2.00855369e+01,  2.98095799e+03])
```

```
In [9]: sigma_e = np.sum(exp)
        sigma_e
```

executed in 9ms, finished 01:47:25 2018-08-21

```
Out[9]: 3553.8854765602264
```

```
In [11]: z = exp/sigma_e
         z
```

executed in 8ms, finished 01:47:34 2018-08-21

```
Out[11]: array([ 1.13517669e-01,  2.81382168e-04,  4.17608165e-02,
                 5.65171192e-03,  8.38788421e-01])
```

Multi Class Classification using MLP

Loss:

- MSE (Mean square error)
- Let predicted label be y' .
- Remains the same even for regression.

Our objective:

- Minimize the difference between y'_i and y_i for all i

$$L(W) = \sum_i ||y'_i - y_i|| \sum_i \sum_j (y'_{ij} - y_{ij})^2$$

Four Cases

Perceptron

$$\mathbf{w}^T \mathbf{x}$$

Linear Features
Linear Classifiers

$$\phi(\mathbf{w})^T \phi(\mathbf{x})$$

Nonlinear Features
Linear Classifiers

Kernel-SVM

MLP

$$\psi(\mathbf{w}, \mathbf{x})$$

Linear Features
Nonlinear Classifiers

$$\psi(\mathbf{w}, \phi(\mathbf{x}))$$

Nonlinear Features
Nonlinear Classifiers

MLP of VGG
Features

Summary

- Many “perceptron” networks can be stacked to generate Multi Layer Perceptron (MLP).
- Any arbitrary function can be approximated.
 - Given that we can train!! (this could be tricky)
- Classically the nonlinearity is a simple sigmoid or similar functions.
- Often people use MLP with one or two hidden layers
 - Not very deep.

Thanks!!

Questions?