

# PyTorch

<https://pytorch.org/>

# Open Source - Deep Learning Frameworks



# Why do we need a special Deep Learning Framework ?

KNN from scratch

VS

KNN from Sklearn

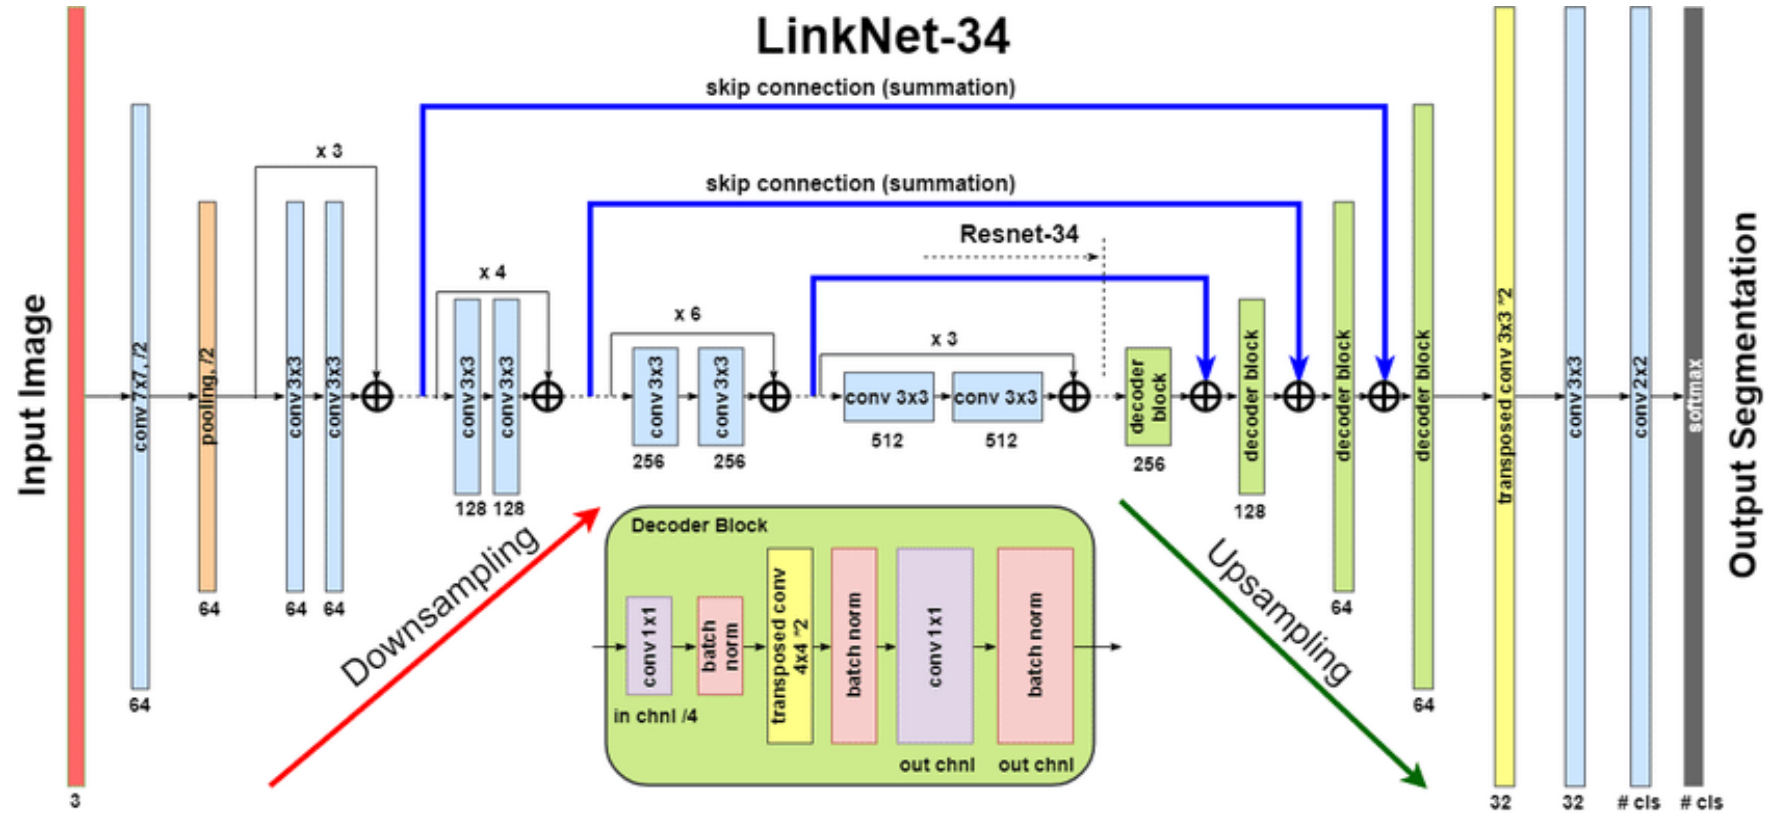
## Python code for Neural Network

# This is just a 2-layered MLP !!

```
def build_model(nn_hdim, num_passes=20000, print_loss=False):  
    np.random.seed(0)  
    W1 = np.random.randn(nn_input_dim, nn_hdim) / np.sqrt(nn_input_dim)  
    b1 = np.zeros((1, nn_hdim))  
    W2 = np.random.randn(nn_hdim, nn_output_dim) / np.sqrt(nn_hdim)  
    b2 = np.zeros((1, nn_output_dim))  
  
    # This is what we return at the end  
    model = {}  
  
    # Gradient descent. For each batch...  
    for i in range(0, num_passes):  
  
        # Forward propagation  
        z1 = X.dot(W1) + b1  
        a1 = np.tanh(z1)  
        z2 = a1.dot(W2) + b2  
        exp_scores = np.exp(z2)  
        probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)  
  
        # Backpropagation  
        delta3 = probs  
        delta3[range(num_examples), y] -= 1  
        dW2 = (a1.T).dot(delta3)  
        db2 = np.sum(delta3, axis=0, keepdims=True)  
        delta2 = delta3.dot(W2.T) * (1 - np.power(a1, 2))  
        dW1 = np.dot(X.T, delta2)  
        db1 = np.sum(delta2, axis=0)  
  
        # Add regularization terms (b1 and b2 don't have regularization terms)  
        dW2 += reg_lambda * W2  
        dW1 += reg_lambda * W1  
  
        # Gradient descent parameter update  
        W1 += -epsilon * dW1  
        b1 += -epsilon * db1  
        W2 += -epsilon * dW2  
        b2 += -epsilon * db2  
  
        # Assign new parameters to the model  
        model = { 'W1': W1, 'b1': b1, 'W2': W2, 'b2': b2}  
  
        # Optionally print the loss.  
        # This is expensive because it uses the whole dataset, so we don't want to do it too often.  
        if print_loss and i % 1000 == 0:  
            print("Loss after iteration %i: %f" % (i, calculate_loss(model)))  
  
    return model
```

```
# Helper function to evaluate the total loss on the dataset  
def calculate_loss(model):  
    W1, b1, W2, b2 = model['W1'], model['b1'], model['W2'], model['b2']  
    # Forward propagation to calculate our predictions  
    z1 = X.dot(W1) + b1  
    a1 = np.tanh(z1)  
    z2 = a1.dot(W2) + b2  
    exp_scores = np.exp(z2)  
    probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)  
    # Calculating the loss  
    correct_logprobs = -np.log(probs[range(num_examples), y])  
    data_loss = np.sum(correct_logprobs)  
    # Add regularization term to loss (optional)  
    data_loss += reg_lambda/2 * (np.sum(np.square(W1)) + np.sum(np.square(W2)))  
    return 1./num_examples * data_loss  
  
# Helper function to predict an output (0 or 1)  
def predict(model, x):  
    W1, b1, W2, b2 = model['W1'], model['b1'], model['W2'], model['b2']  
    # Forward propagation  
    z1 = x.dot(W1) + b1  
    a1 = np.tanh(z1)  
    z2 = a1.dot(W2) + b2  
    exp_scores = np.exp(z2)  
    probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)  
    return np.argmax(probs, axis=1)
```

# What about coding for this network



And finding the gradients !!

# Standard Deep Learning Framework

## Computational graphs


Numerical Examples

$$f(x, y, z) = \sigma(w_0 x_0 + w_1 x_1 + b)$$

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

$$\frac{d\sigma(x)}{dx} = \frac{e^{-x}}{(1 + e^{-x})^2} = \left( \frac{1 + e^{-x} - 1}{(1 + e^{-x})} \right) \left( \frac{1}{(1 + e^{-x})} \right)$$

$$\frac{d\sigma(x)}{dx} = (1 - \sigma(x))\sigma(x)$$




From Stanford Course: Convolutional Neural Networks for Visual Recognition

# What is computational graph?

- A directed graph, where every node represents a mathematical operation and edges represent the variables which are to be operated.

$$f(x, y, z) = (x + y)z$$

$$eg : x = -2, y = 5, z = -4$$



# Types of Computational Graphs

## Static

Cannot add nodes at runtime

Building graph is separated  
from execution

## Dynamic

Allows addition of nodes at  
runtime

Every execution there is a  
new graph



# Standard Deep Learning Framework

**Build and operate  
Computational Graphs**

**Auto-differentiation**

Compute and take derivatives of  
huge composition functions


# Standard Deep Learning Framework

**Build and operate  
Computational Graphs**

**Auto-differentiation**

**Parallelizing on GPU**

# Parallelizing



GPU  
 Many cores

# Standard Deep Learning Framework

**Build and operate  
Computational Graphs**

**Perform forward  
and backward  
propagation**

**Parallelizing on GPU**

**Provide with Standard  
Architectures and other widely  
used primitives**

# Pytorch

**FAIR (Facebook Artificial  
Intelligence Research)**

**Based on Torch**

Dynamic  
Computational  
graph

Easy to  
implement,  
debug,  
Developer  
Friendly


Tensor  
computation  
with strong  
GPU  
acceleration

Efficient  
Memory Usage


# Basics of Pytorch

# The building blocks

## Tensors



## Variables



- A Variable class wraps a tensor. You can access this tensor by calling `.data` attribute of a Variable.
- The Variable also stores the gradient of a scalar quantity (say, loss) with respect to the parameter it holds. This gradient can be accessed by calling the `.grad` attribute.
- The third attribute a Variable holds is a `grad_fn`, a Function object which created the variable.
  - $c = a + b$ . Then `c` is a new variable, and its `grad_fn` is something called `AddBackward`

# The building blocks

Tensors

Like Numpy arrays but can run on GPU

GPU

```
import torch

dtype = torch.FloatTensor

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)
w1 = torch.randn(D_in, H).type(dtype)
w2 = torch.randn(H, D_out).type(dtype)
```

```
import torch

dtype = torch.cuda.FloatTensor

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)
w1 = torch.randn(D_in, H).type(dtype)
w2 = torch.randn(H, D_out).type(dtype)
```



# The building blocks

Variables are nodes in a computational graph which store the data and gradient

Variables

```
import torch
from torch.autograd import Variable

N, D_in, H, D_out = 64, 1000, 100, 10
x = Variable(torch.randn(N, D_in), requires_grad=False)
y = Variable(torch.randn(N, D_out), requires_grad=False)
w1 = Variable(torch.randn(D_in, H), requires_grad=True)
w2 = Variable(torch.randn(H, D_out), requires_grad=True)
```

# Difference

```
import torch

dtype = torch.FloatTensor

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)
w1 = torch.randn(D_in, H).type(dtype)
w2 = torch.randn(H, D_out).type(dtype)
```

```
learning_rate = 1e-6
for t in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()
```

```
grad_y_pred = 2.0 * (y_pred - y)
grad_w2 = h_relu.t().mm(grad_y_pred)
grad_h_relu = grad_y_pred.mm(w2.t())
grad_h = grad_h_relu.clone()
grad_h[h < 0] = 0
grad_w1 = x.t().mm(grad_h)
```

```
w1 -= learning_rate * grad_w1
w2 -= learning_rate * grad_w2
```

```
import torch
from torch.autograd import Variable
N, D_in, H, D_out=64,1000,100,10
x=Variable(torch.randn(N, D_in), requires_grad=False)
y=Variable(torch.randn(N, D_out), requires_grad=False)
w1=Variable(torch.randn(D_in,H), requires_grad=True)
w2=Variable(torch.randn(H, D_out), requires_grad=True)
learning_rate=1e-6
for t in range(500):
    y_pred=x.mm(w1).clamp(min=0).mm(w2)
    loss=(y_pred-y).pow(2).sum()
    loss.backward()
    w1.data-=learning_rate*w1.grad
    w2.data-=learning_rate*w2.grad
    w1.grad.data.zero_()
    w2.grad.data.zero_()
```

# Neural Network

**torch.nn**

# NN module

Higher-level wrappers for working with neural nets

- Layers
- Activation Functions
- Loss functions

torch.nn

Parameters

Containers

Convolution Layers

Conv1d

Conv2d

Conv3d

ConvTranspose1d

ConvTranspose2d

ConvTranspose3d

Other layers:  
Dropout, Linear,  
Normalization Layer

torch.nn

Parameters

Containers

Convolution Layers

Pooling Layers

MaxPool1d

MaxPool2d

MaxPool3d

MaxUnpool1d

MaxUnpool2d

MaxUnpool3d

AvgPool1d

AvgPool2d

AvgPool3d

FractionalMaxPool2d

LPPool2d

AdaptiveMaxPool1d

AdaptiveMaxPool2d

AdaptiveMaxPool3d

AdaptiveAvgPool1d

AdaptiveAvgPool2d

AdaptiveAvgPool3d

Loss functions

L1Loss

MSELoss

CrossEntropyLoss

NLLLoss

PoissonNLLLoss

KLDivLoss

BCELoss

BCEWithLogitsLoss

MarginRankingLoss

HingeEmbeddingLoss

MultiLabelMarginLoss

SmoothL1Loss

SoftMarginLoss


MultiLabelSoftMarginLoss

CosineEmbeddingLoss

MultiMarginLoss

TripletMarginLoss

# NN Block Diagram



```


class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(320, 10) # 320 -> 10

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = F.relu(self.mp(self.conv2(x)))
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
  
```

torch.nn.Conv2d(in\_channels, out\_channels, kernel\_size, ...)  
 torch.nn.MaxPool2d(kernel\_size, ..)

# NN Block Diagram



```

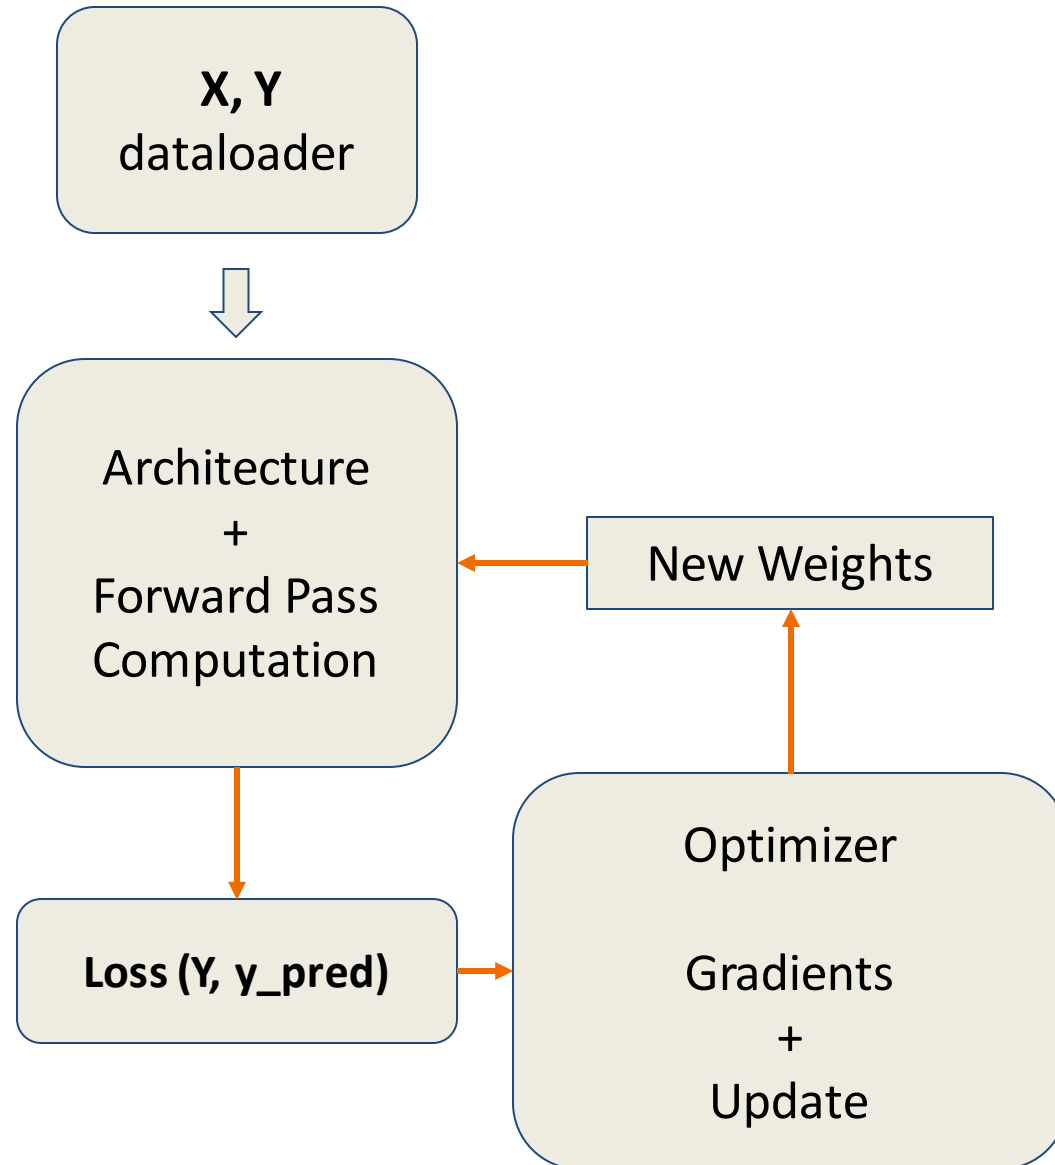
import torch
import torch.nn as nn
class Net(nn.Module):
    def __init__(self, input_size, hidden_size, num_classes):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_size, num_classes)
        self.softmax = nn.Softmax(dim=1)

        #Weight Initialization
        for m in self.modules():
            if isinstance(m, nn.Linear):
                weight_init.xavier_normal_(m.weight)

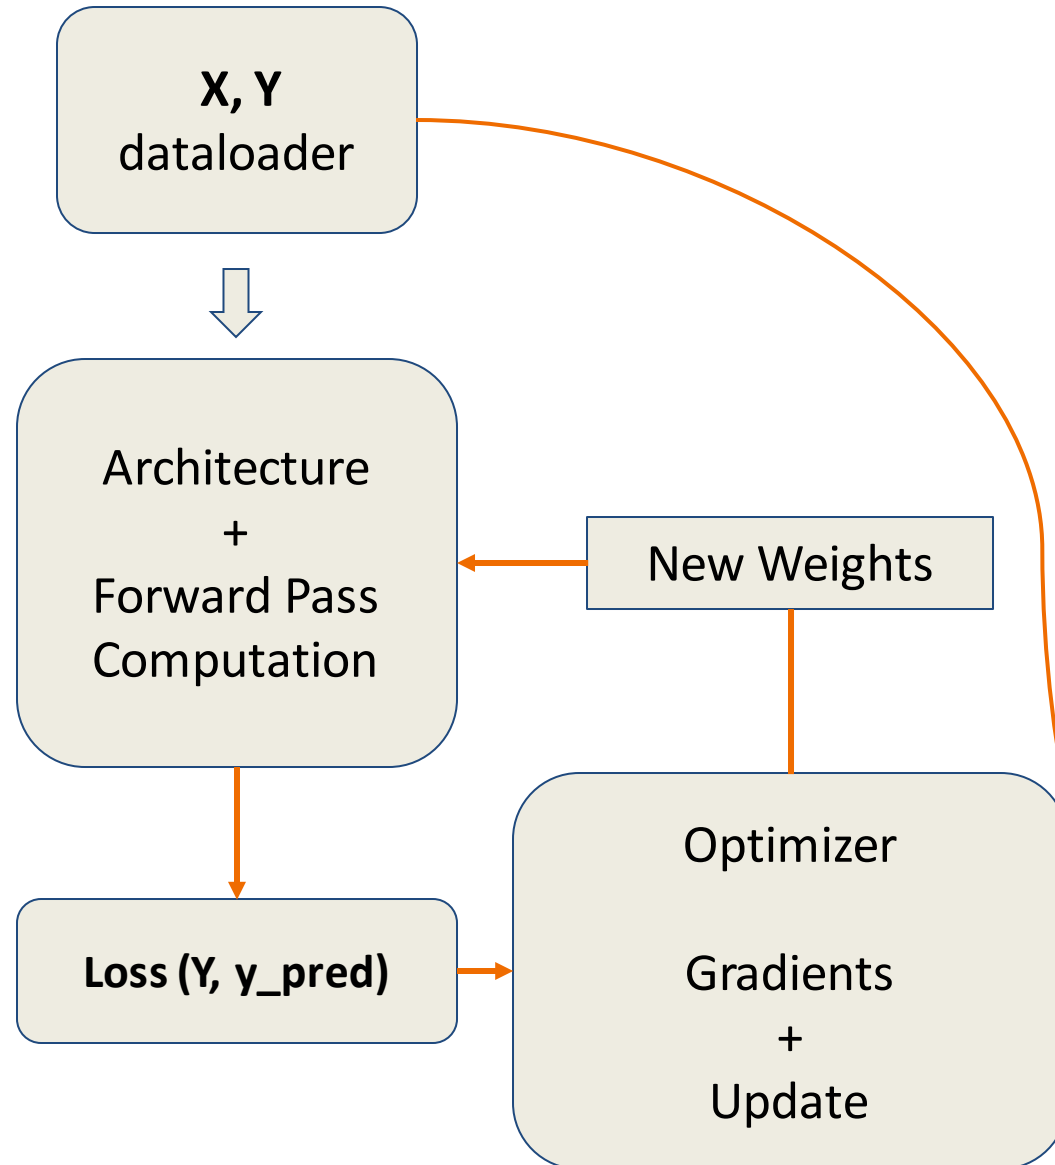
    def forward(self, x):
        out = self.fc1(x)
        out = self.relu(out)
        out = self.fc2(out)
        out = self.softmax(out)
        return out

N, D_in, H, D_out = 64, 1000, 100, 10
net = Net(D_in, H, D_out)
criterion = nn.CrossEntropyLoss()
  
```

# NN Training



# NN Training



```

import torch
import torch.nn as nn
class Net(nn.Module):
    def __init__(self, input_size, hidden_size, num_classes):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_size, num_classes)
        self.softmax = nn.Softmax(dim=1)

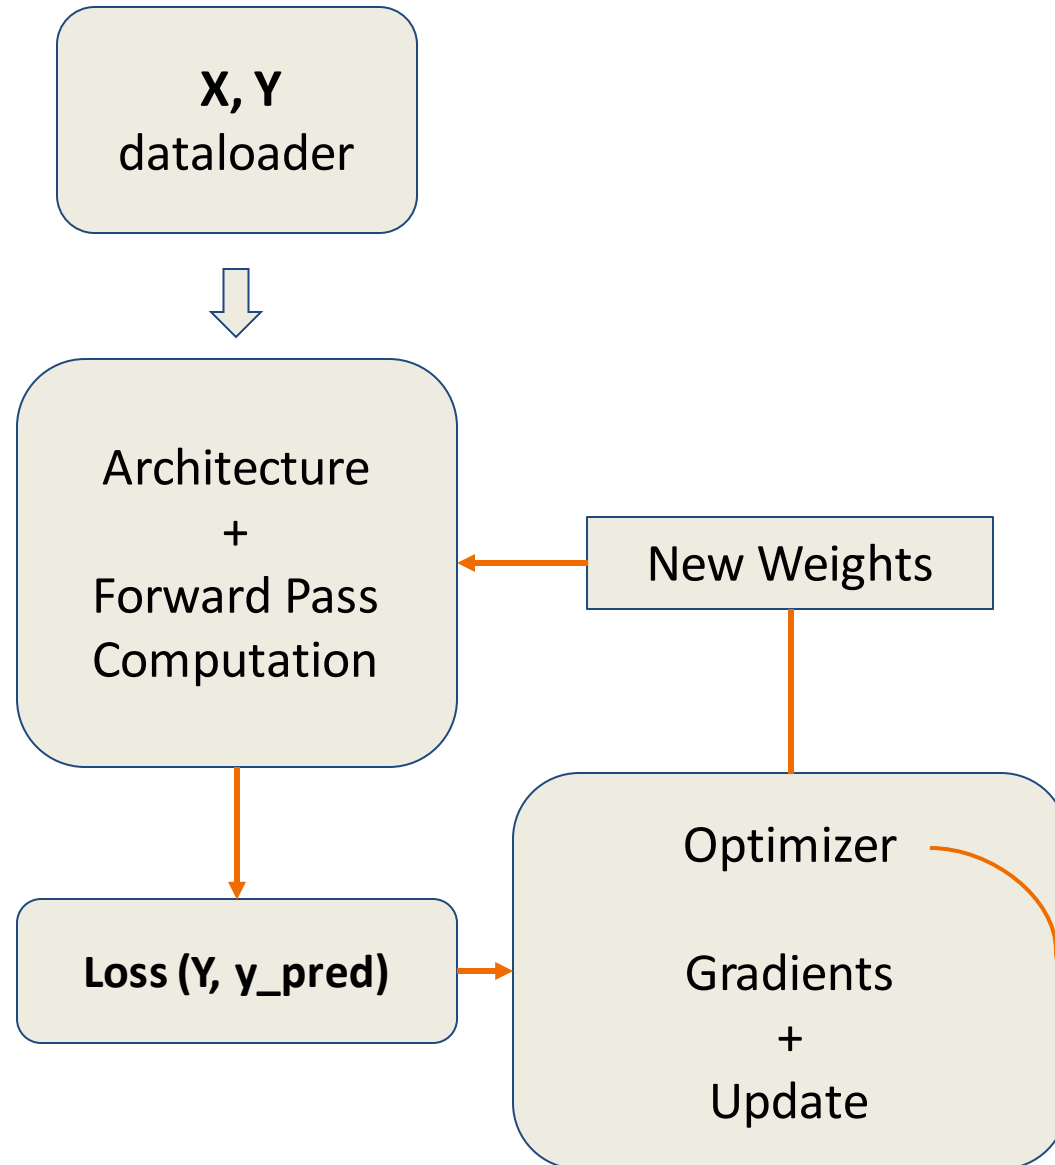
        #Weight Initialization
        for m in self.modules():
            if isinstance(m, nn.Linear):
                weight_init.xavier_normal_(m.weight)

    def forward(self, x):
        out = self.fc1(x)
        out = self.relu(out)
        out = self.fc2(out)
        out = self.softmax(out)
        return out

N, D_in, H, D_out = 64, 1000, 100, 10
net=Net(D_in, H, D_out)
criterion=nn.CrossEntropyLoss()
X=torch.randn(N, D_in)
y=torch.randn(N, D_out)
train_loader=Dataloader(TensorDataset(X,y), batchsize=8)
  
```



# NN Training



```

use_cuda = torch.cuda.is_available()
device = torch.device("cuda" if use_cuda else "cpu")

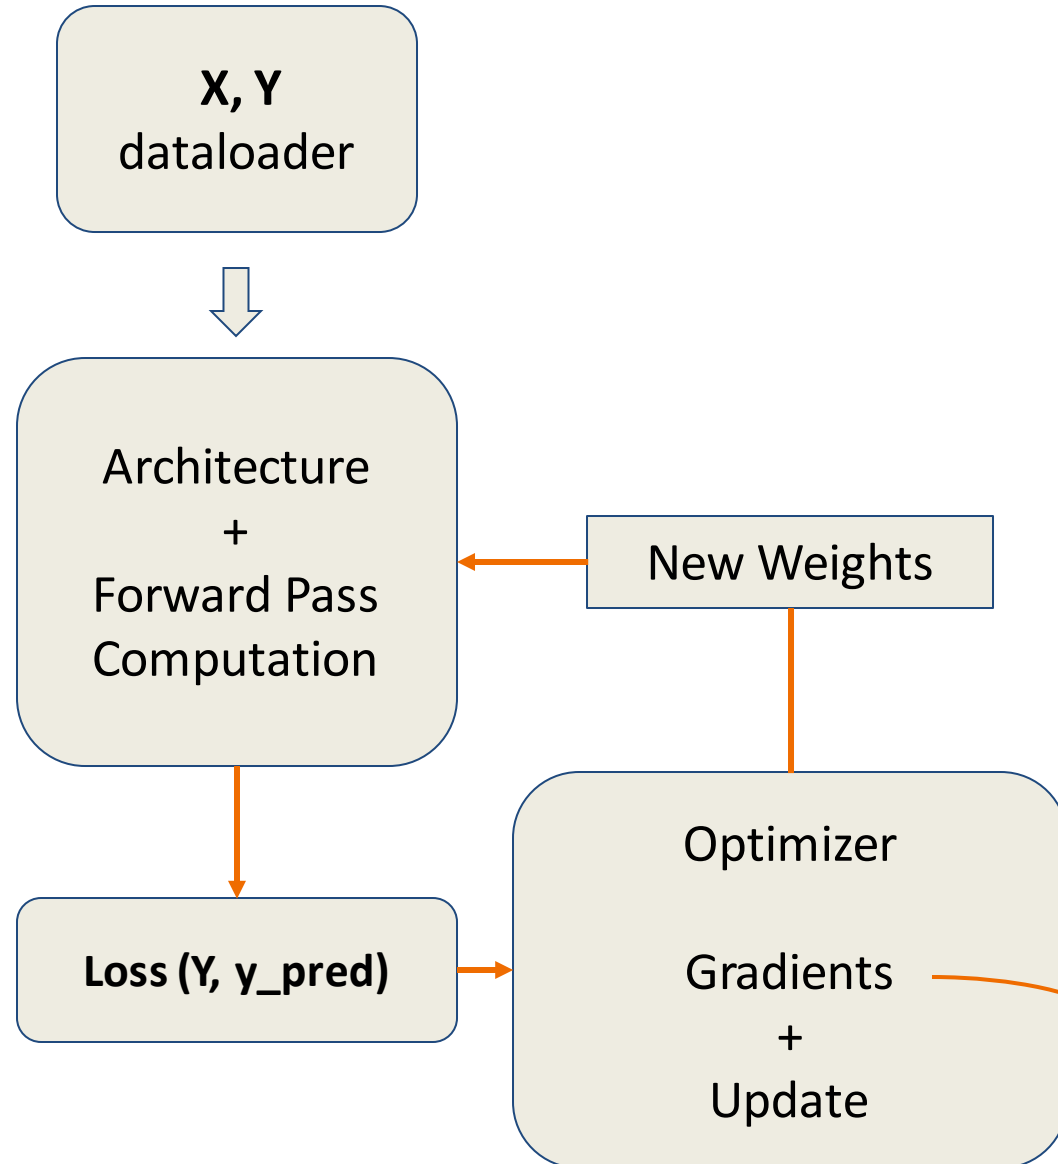
import torch
import torch.nn as nn
class Net(nn.Module):
    def __init__(self, input_size, hidden_size, num_classes):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_size, num_classes)
        self.softmax = nn.Softmax(dim=1)

        #Weight Initialization
        for m in self.modules():
            if isinstance(m, nn.Linear):
                weight_init.xavier_normal_(m.weight)

    def forward(self, x):
        out = self.fc1(x)
        out = self.relu(out)
        out = self.fc2(out)
        out = self.softmax(out)
        return out

N, D_in, H, D_out = 64, 1000, 100, 10
net=Net(D_in, H, D_out)
criterion=nn.CrossEntropyLoss()
X=torch.randn(N, D_in)
y=torch.randn(N, D_out)
train_loader=Dataloader(TensorDataset(X,y), batchsize=8)
optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate)
net = net.to(device)
  
```

# NN Training



```

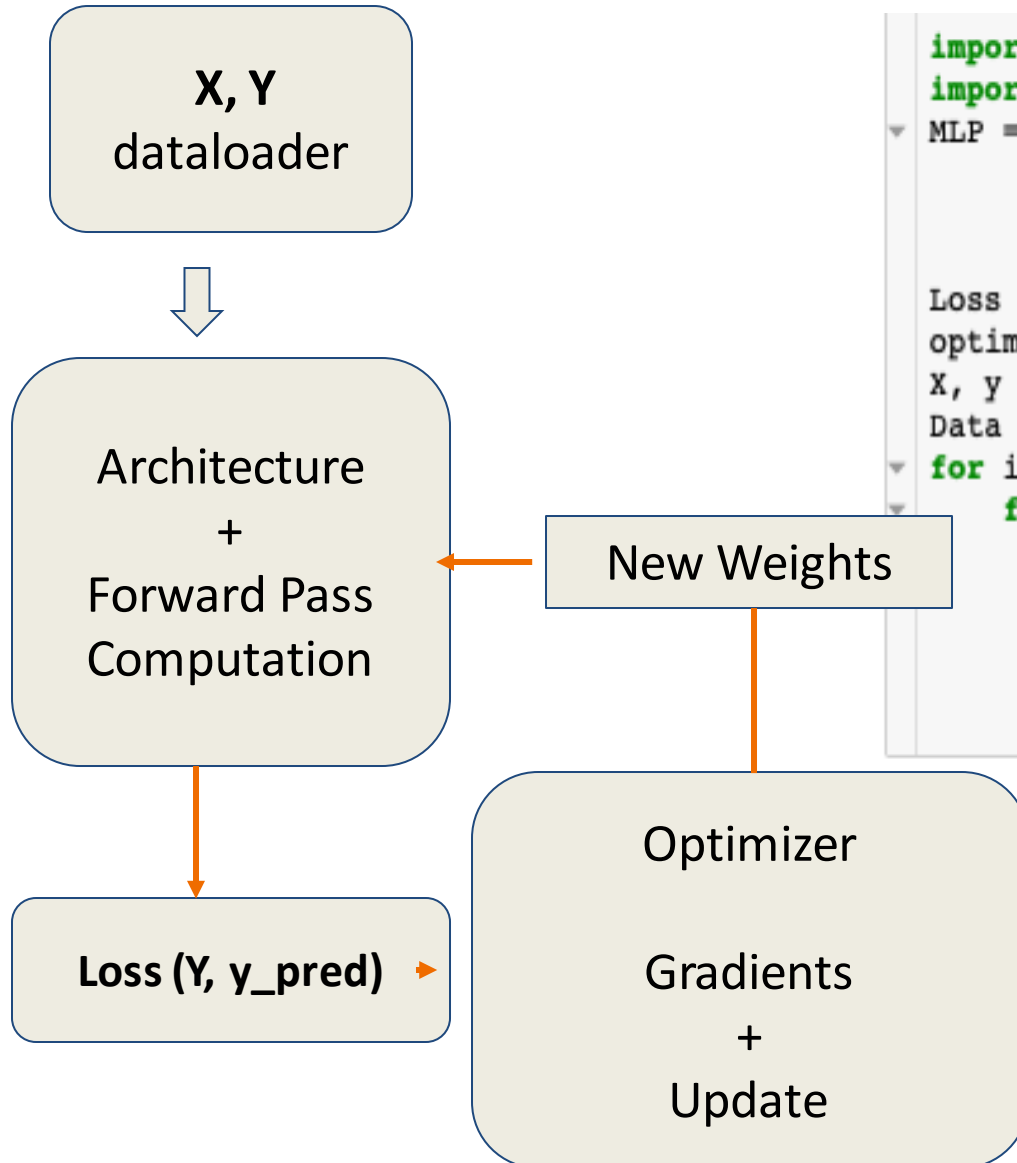
class Net(nn.Module):
    def __init__(self, input_size, hidden_size, num_classes):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_size, num_classes)
        self.softmax = nn.Softmax(dim=1)

        #Weight Initialization
        for m in self.modules():
            if isinstance(m, nn.Linear):
                weight_init.xavier_normal_(m.weight)

    def forward(self, x):
        out = self.fc1(x)
        out = self.relu(out)
        out = self.fc2(out)
        out = self.softmax(out)
        return out

N, D_in, H, D_out = 64, 1000, 100, 10
net=Net(D_in, H, D_out)
criterion=nn.CrossEntropyLoss()
X=torch.randn(N, D_in)
y=torch.randn(N, D_out)
train_loader=Dataloader(TensorDataset(X,y), batchsize=8)
optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate)
net = net.to(device)
for epoch in range(num_epochs):
    for i, (images, labels) in enumerate(train_loader):
        images = images.view(-1, 28*28)
        labels = labels
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = net(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
  
```

# NN Training (alternate)

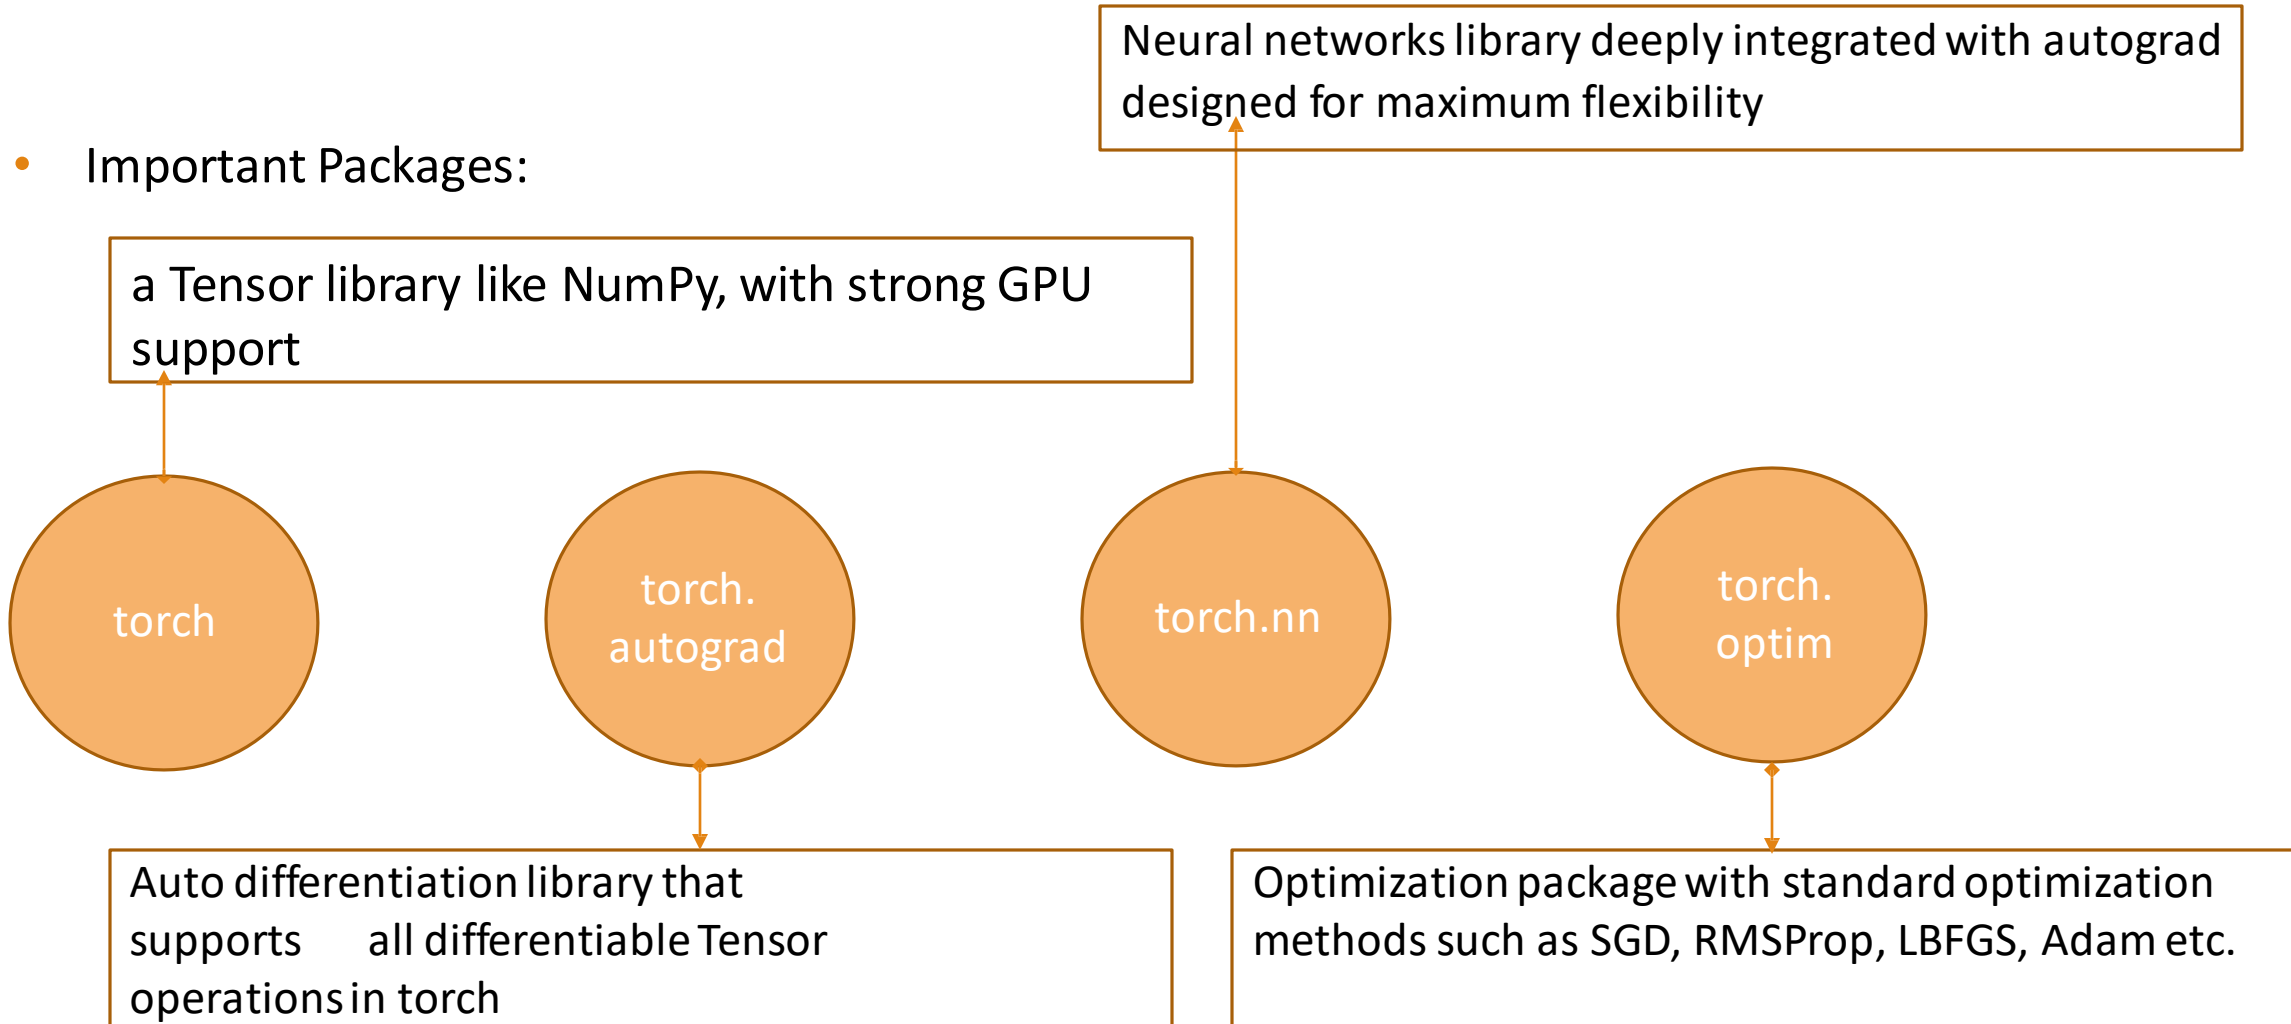


```

import torch
import torch.nn as nn
MLP = nn.Sequential(nn.Linear(2,3)
                    ,nn.ReLU()
                    , nn.Linear(3,2)
                    ,nn.Softmax()).double()
Loss = nn.CrossEntropyLoss()
optimiser = torch.optim.SGD(MLP.parameters(), lr = 0.01, weight_decay= 0.01)
X, y = sklearn.datasets.make_moons(200, noise=0.20)
Data = zip(torch.tensor(X), torch.tensor(y))
for i in range(10):
    for x,y in Data:
        optimiser.zero_grad()
        ypred = MLP(x)
        loss = Loss(ypred.unsqueeze(0),y.unsqueeze(0))
        loss.backward()
        optimiser.step()
  
```

# PyTorch

- Important Packages:



---

# Computation on GPU

---



# Computational Graph: GPU

```
class Net(nn.Module):
    def __init__(self, input_size, hidden_size, num_classes):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_size, num_classes)
        self.softmax = nn.Softmax(dim=1)

        #Weight Initialization
        for m in self.modules():
            if isinstance(m, nn.Linear):
                weight_init.xavier_normal_(m.weight)

    def forward(self, x):
        out = self.fc1(x)
        out = self.relu(out)
        out = self.fc2(out)
        out = self.softmax(out)
        return out

N, D_in, H, D_out = 64, 1000, 100, 10
net = Net(D_in, H, D_out)
criterion = nn.CrossEntropyLoss()
X = torch.randn(N, D_in)
y = torch.randn(N, D_out)
train_loader = DataLoader(TensorDataset(X, y), batchsize=8)
optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate)
net = net.to(device)
for epoch in range(num_epochs):
    for i, (images, labels) in enumerate(train_loader):
        images = images.view(-1, 28*28)
        labels = labels
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = net(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
```



# Computational Graph: Pytorch

## Computational Graph with Pytorch on GPU

```

: import torch

: use_cuda = torch.cuda.is_available()
: device = torch.device("cuda" if use_cuda else "cpu")

: x=torch.randn(3,4, requires_grad=True).to(device)
: y=torch.randn(3,4, requires_grad=True).to(device)
: z=torch.randn(3,4, requires_grad=True).to(device)

: a = x*y
: b = a+z
: c = torch.sum(b)


: c.backward()

: print(x.grad.data, '\n', y.grad.data, z.grad.data)

```

tensor([[ 2.0230, 1.5964, 0.2990, 1.0220],  
 [-0.3020, -1.5197, -1.1080, -0.3008],  
 [-1.9656, -0.7705, -1.2423, -0.0137]]) tensor([[ -1.8837, -1.5356, -0.4437, 1.4743],  
 [-0.7601, -1.1921, -0.6407, 0.6169],  
 [ 1.5937, -0.0240, 0.5942, -0.7387]]) tensor([[1., 1., 1., 1.],  
 [1., 1., 1., 1.],  
 [1., 1., 1., 1.]])

Computation happens on GPU, if available



# CPU vs GPU

## CPU vs GPU

	# Cores	Clock Speed	Memory	Price
<b>CPU</b> (Intel Core i7-7700k)	4 (8 threads with hyperthreading)	4.4 GHz	Shared with system	\$339
<b>CPU</b> (Intel Core i7-6950X)	10 (20 threads with hyperthreading)	3.5 GHz	Shared with system	\$1723
<b>GPU</b> (NVIDIA Titan Xp)	3840	1.6 GHz	12 GB GDDR5X	\$1200
<b>GPU</b> (NVIDIA GTX 1070)	1920	1.68 GHz	8 GB GDDR5	\$399

**CPU:** Fewer cores, but each core is much faster and much more capable; great at sequential tasks

**GPU:** More cores, but each core is much slower and “dumber”; great for parallel tasks



**Thanks!!**

**Questions?**