

Self-Attention Layer




What happens in a layer ?

- Input: a set/sequence/grid of vectors of representations
- Output: a set/sequence/grid of vectors of representation
- **Solution:**
 - RNN Layer
 - CNN Layer
 - Self Attention Layer/Transformer
- **Deep Neural Architecture**
 - Representation goes through a sequence of I/O transformations that enrich the semantics and make it more suitable for tasks

Conceptual Comparison

Recurrent Neural Network




Works on **Ordered Sequences**

(+) **Good at long sequences:** After one RNN layer, h_T "sees" the whole sequence

(-) **Not parallelizable:** need to compute hidden states sequentially

1D Convolution

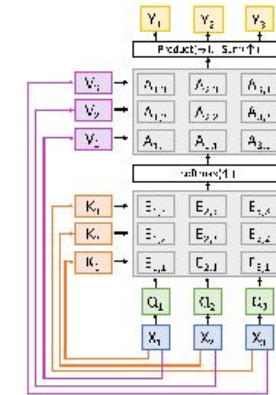


Works on **Multidimensional Grids**

(-) **Bad at long sequences:** Need to stack many conv layers for outputs to "see" the whole sequence

(+) **Highly parallel:** Each output can be computed in parallel

Self-Attention



Works on **Sets of Vectors**

(-) **Good at long sequences:** after one self-attention layer, each output "sees" all inputs!

(+) **Highly parallel:** Each output can be computed in parallel

(-) **Very memory intensive**

Convolution Layer Vs SA Layer

Blank Slide

Comments

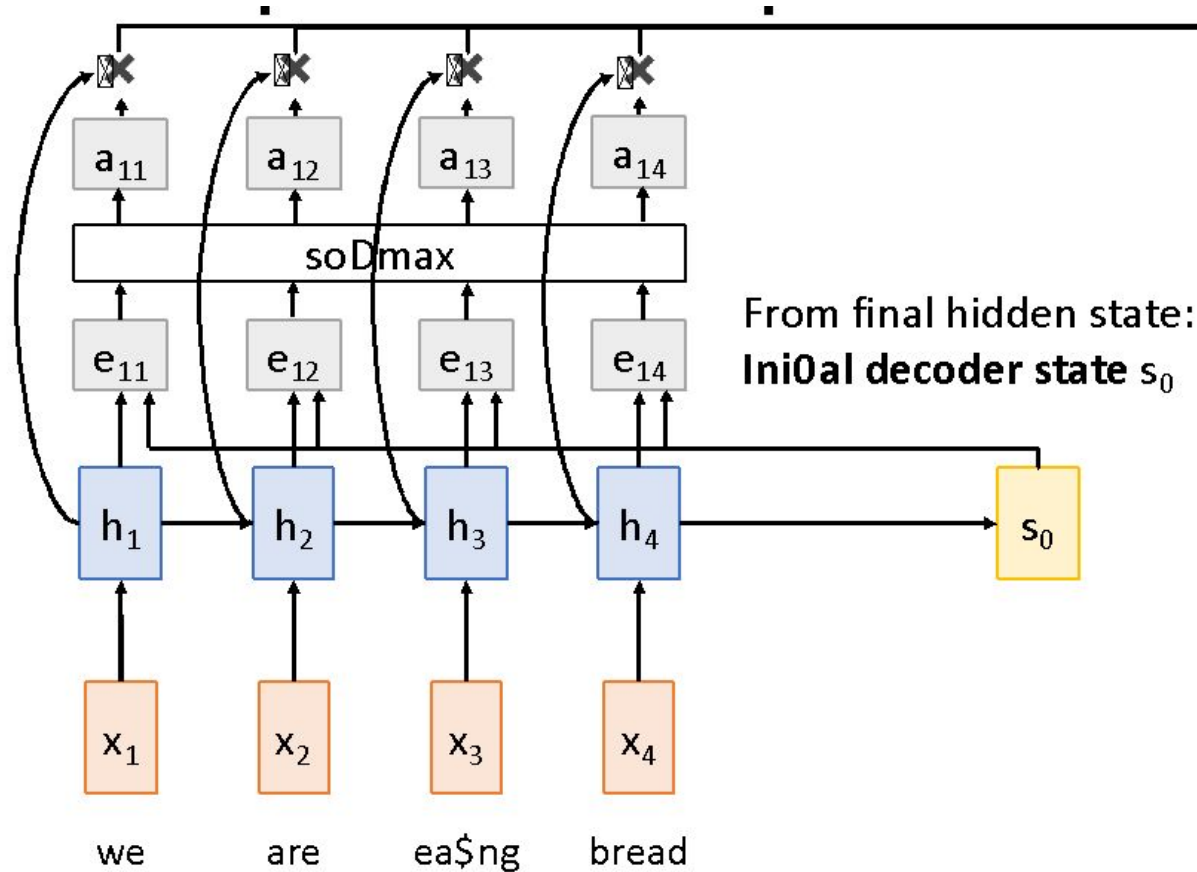
Challenges with RNNs

- Long range dependencies
- Gradient vanishing and explosion
- Large # of training steps
- Recurrence prevents parallel computation

Transformer Networks

- Facilitate long range dependencies
- No gradient vanishing and explosion
- Fewer training steps
- No recurrence that facilitate parallel computation

Attention in RNNs



Compute (scalar) **alignment scores**
 $e_{t,i} = f_{aE}(s_{t-1}, h_i)$ (f_{aE} is an MLP)

Normalize alignment scores
 to get **attention weights**
 $0 < a_{t,i} < 1 \quad \sum_i a_{t,i} = 1$

Compute context vector as linear
 combination of hidden states
 $c_t = \sum_i a_{t,i} h_i$


Use context vector in
 decoder: $s_t = g_U(y_{t-1}, s_{t-1}, c_t)$

**This is all differentiable! Do not
 supervise attention weights –
 backprop through everything**

Attention in RNNs (Seq2Seq)

Use a different context vector in each timestep of decoder

- Input sequence not bottlenecked through single vector
- At each timestep of decoder, context vector “looks at” different parts of the input sequence



Blank Slide

Attention in CNNs (Image Captioning)

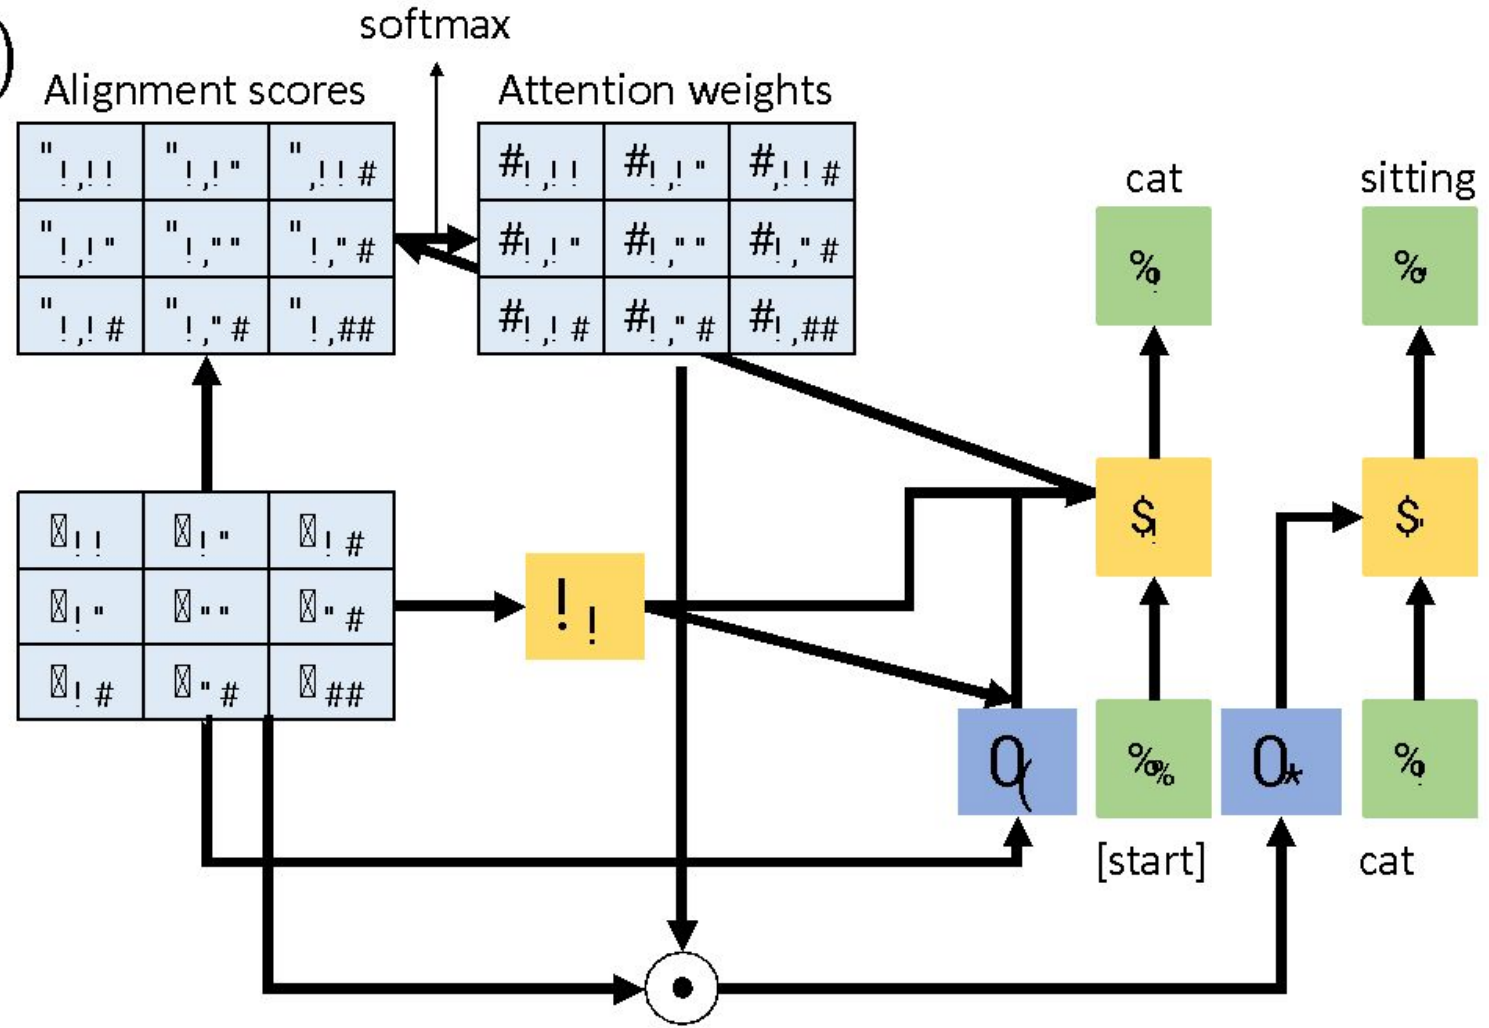
$$a_{ij} = \frac{\exp(s_{ij})}{\sum_k \exp(s_{ik})}$$

$$s_{ij} = \mathbf{v}_i^T \mathbf{W} \mathbf{f}_j$$

$$O_i = \sum_j a_{ij} \mathbf{f}_j$$



CNN



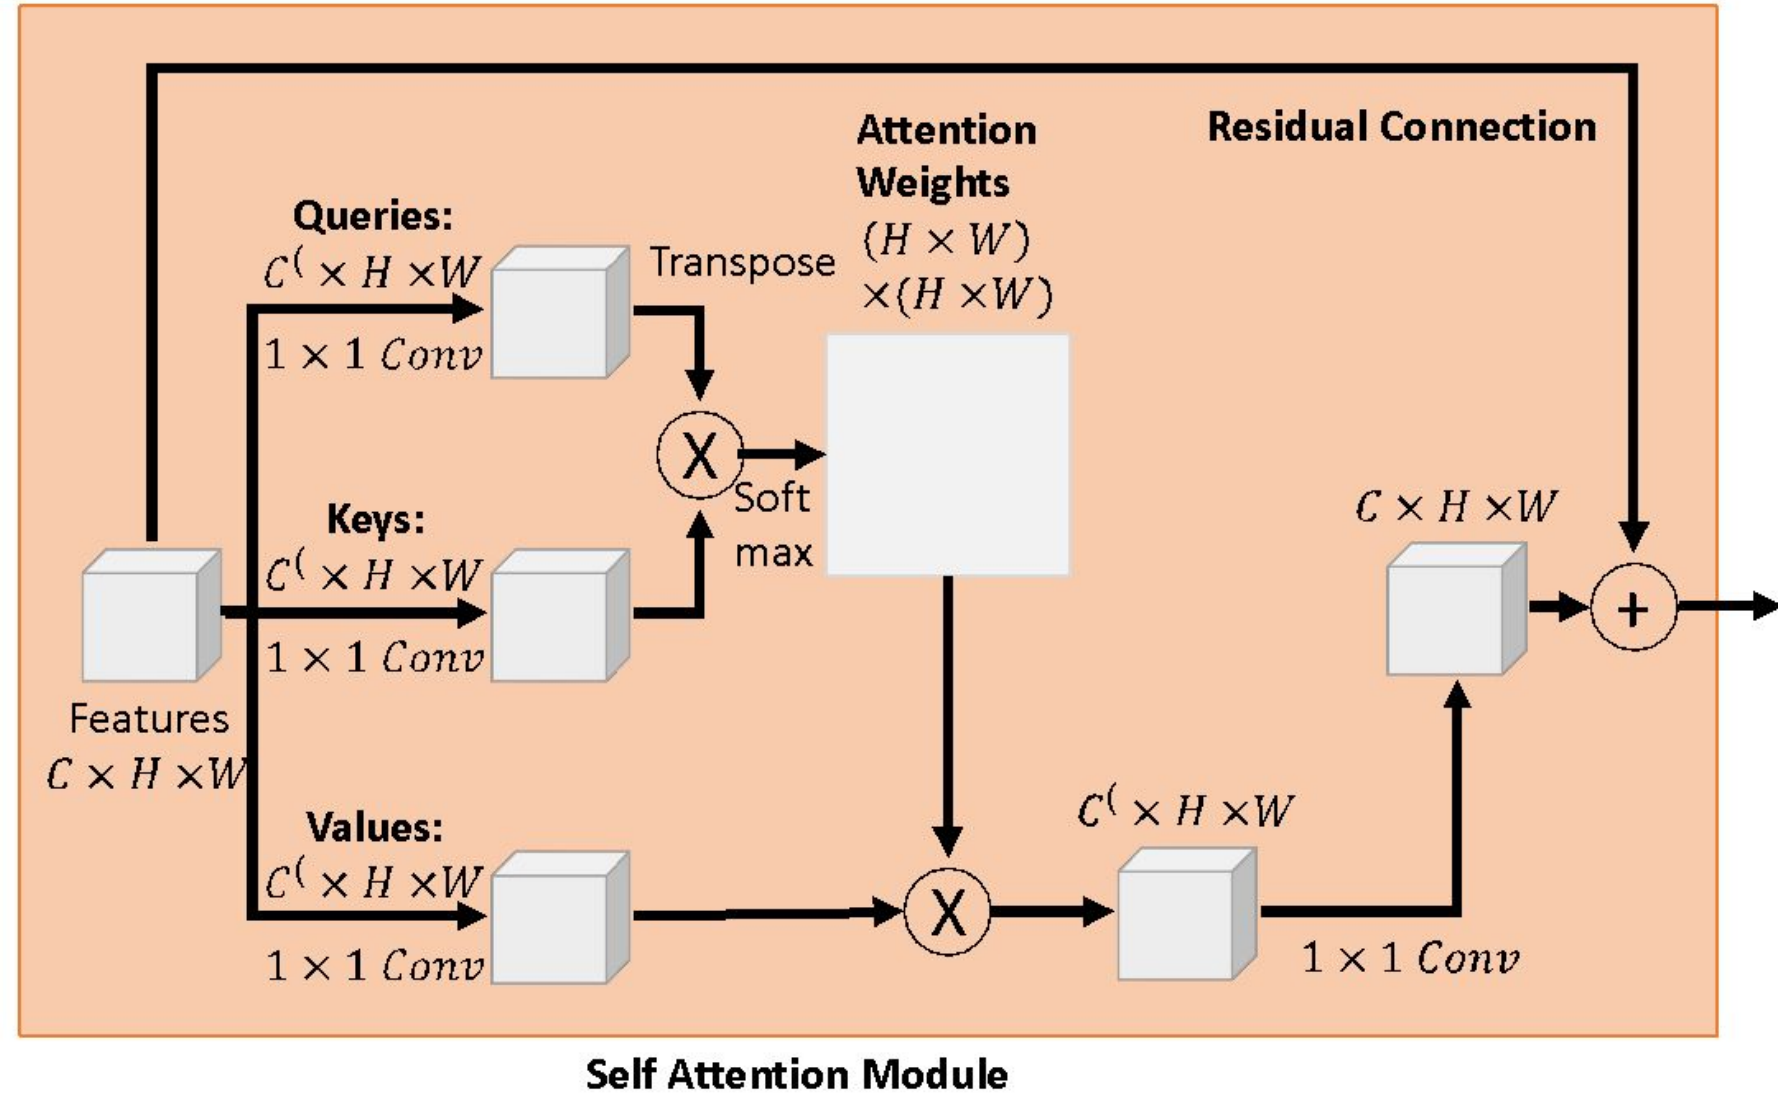
Use a CNN to compute a grid of features for an image

Attention in CNNs

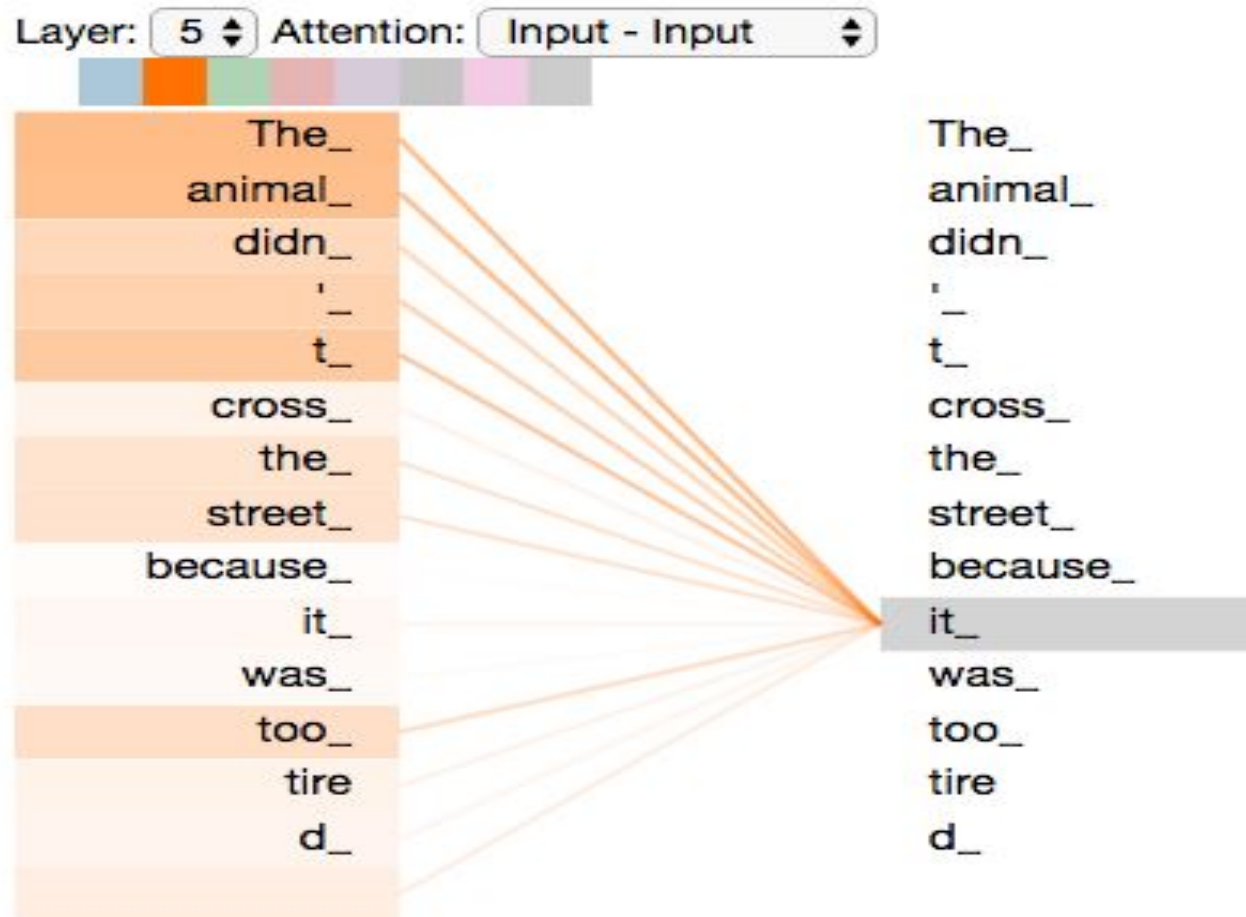
Input Image



CNN



Self Attention



1. **The animal** did not cross the road because **it** was too tired.
2. The animal did not cross **the road** because **it** was too wide.

What do we want to happen in the SA layer?
How does it compare with what we do in FC, or CNN?

Blank Slide

Self Attention: Comments

- Self attention learns the relationship between elements in a sequence.
 - say between words in a sentence
- Self Attention Vs Convolution
 - Filters are dynamically calculated instead of static filters
 - SA is invariant to changes in the input points
 - SA can operate on irregular inputs
- SA allows to learn global and local features
 - Hierarchical feature learning by cascading

Query, Key and Value

We project each embedding: **Queries** **Keys** **Values**

Queries: "Here's what I'm looking for" $W^Q \in \mathbb{R}^{D \times d_k}$

Keys: "Here's what I have" $W^K \in \mathbb{R}^{D \times d_k}$

Values: "What gets communicated" $W^V \in \mathbb{R}^{D \times d_v}$

d_k is dimension of **queries** & **keys**, d_v is dimension of **values**

$$Q = XW^Q \in \mathbb{R}^{N \times d_k} \quad K = XW^K \in \mathbb{R}^{N \times d_k} \quad V = XW^V \in \mathbb{R}^{N \times d_v}$$

Query, Key and Value

- As *the current focus of attention* when being compared to all of the other preceding inputs. We'll refer to this role as a **query**.
- In its role as *a preceding input* being compared to the current focus of attention. We'll refer to this role as a **key**.
- And finally, as a **value** used to compute the output for the current focus of attention.

$$\mathbf{q}_i = \mathbf{x}_i \mathbf{W}^Q; \quad \mathbf{k}_i = \mathbf{x}_i \mathbf{W}^K; \quad \mathbf{v}_i = \mathbf{x}_i \mathbf{W}^V$$

Blank Slide

Computation

$$q_i = x_i W^Q; k_i = x_i W^K; v_i = x_i W^V$$

$$\text{score}(x_i, x_j) = \frac{q_i \cdot k_j}{d_k}$$


$$a_{ij} = \frac{\text{softmax}(\text{score}(x_i, x_j))}{X} \quad \forall j \leq i$$

$$a_i = \sum_{j \leq i} a_{ij} v_j$$

Blank Slide

Conceptual Comparison

Recurrent Neural Network




Works on **Ordered Sequences**

(+) **Good at long sequences:** After one RNN layer, h_T "sees" the whole sequence

(-) **Not parallelizable:** need to compute hidden states sequentially

1D Convolution




Works on **Multidimensional Grids**

(-) **Bad at long sequences:** Need to stack many conv layers for outputs to "see" the whole sequence

(+) **Highly parallel:** Each output can be computed in parallel

Self-Attention



Works on **Sets of Vectors**

(-) **Good at long sequences:** after one self-attention layer, each output "sees" all inputs!

(+) **Highly parallel:** Each output can be computed in parallel

(-) **Very memory intensive**

Convolution Layer Vs SA Layer

Blank Slide

Blank Slide

Blank Slide

Where we want to go?

The Transformer

Transformer Block:


Input: Set of vectors x

Output: Set of vectors y

Self-attention is the only interaction between vectors!

Layer norm and MLP work independently per vector

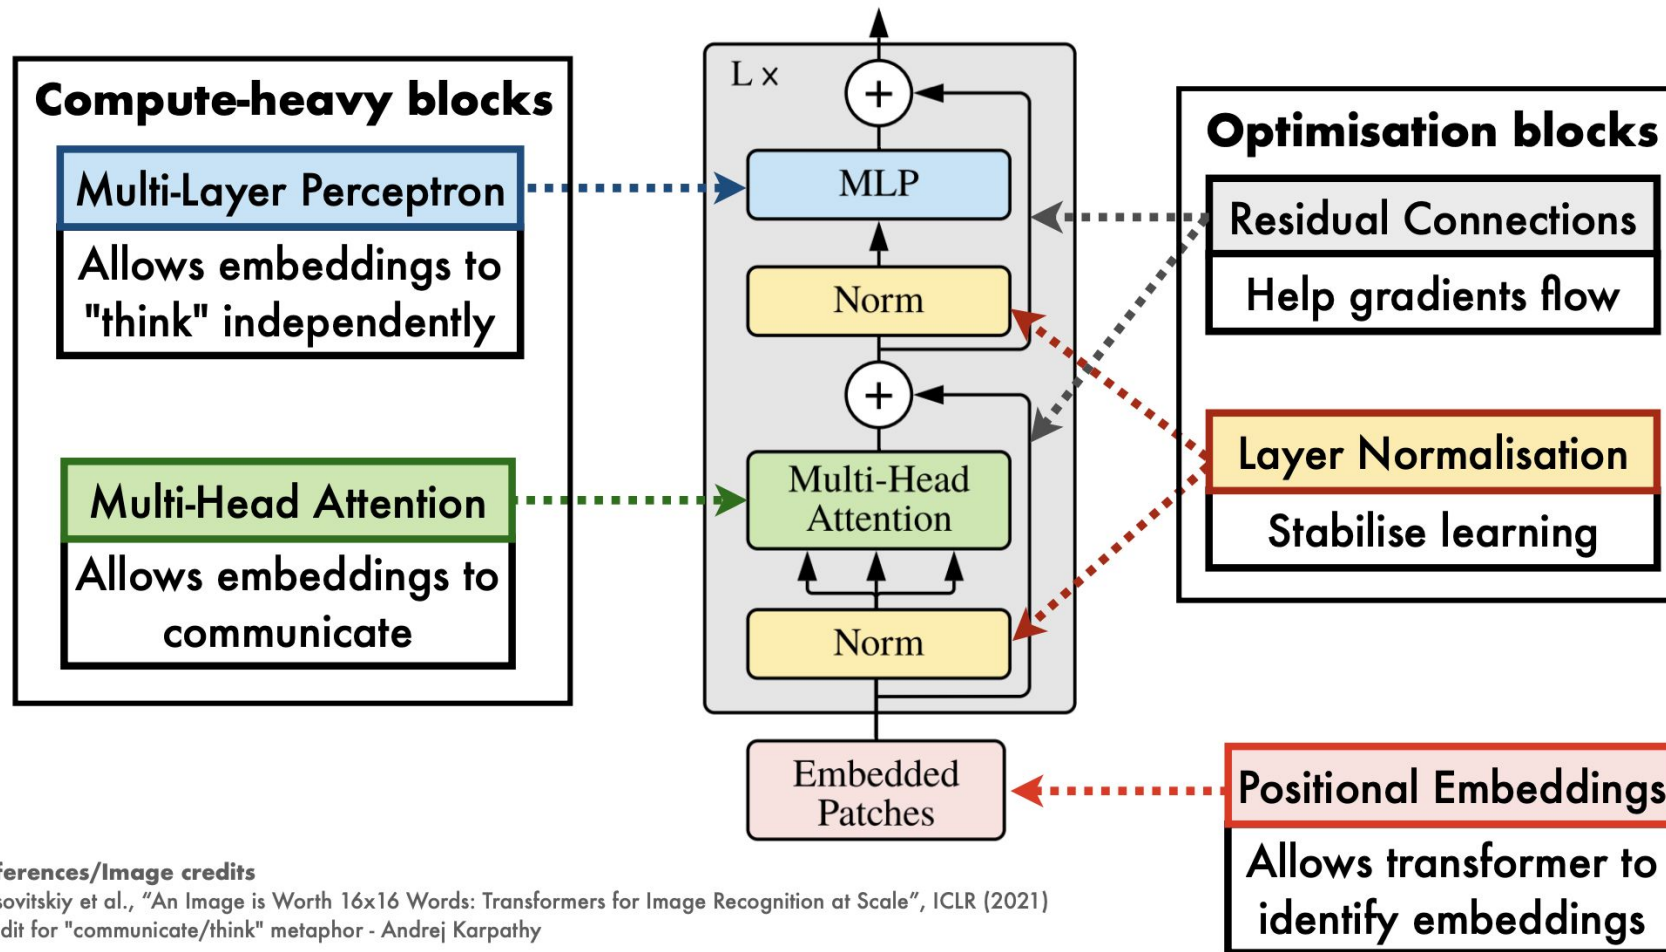
Highly scalable, highly parallelizable



Transformers

Transformer Encoder

Five key ideas



References/Image credits

Dosovitskiy et al., "An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale", ICLR (2021)

Credit for "communicate/think" metaphor - Andrej Karpathy

Blank Slide

Blank Slide

Thanks!!

Questions?